

Waived in America

Procurement Transparency As Industrial Policy

Anantesh Mohapatra

May 2026

An honors thesis submitted in partial fulfillment of the requirements of the degree:

Bachelor of Science
Undergraduate College
Leonard N. Stern School of Business
New York University

Thesis Advisor: Professor Christopher Conlon

Faculty Advisor: Professor Lawrence White

Contents

Acknowledgments	1
Introduction	2
Background	4
The History of Buy American	4
Before Centralization	5
The New Waiver System	6
Literature Review	7
The Cost of Buying American	7
Information Transparency and Sourcing Outcomes	7
The Motivation Behind Public Waivers	8
Waiver Hypotheses	10
Hypothesis 1: Domestic Sourcing	10
Hypothesis 2: Competition	10
Hypothesis 3: Price	10
Research Design	12
Unit of Analysis	12
Treatment Panel	13
Event Study	14
Matched Controls	14
Synthetic Controls	15
Data	17
Sources	17
Descriptive Statistics for Waivers	18
Procurement in Context	21

From Waiver Universe to Analysis Sample	22
Building the Analysis Panel	24
Measuring Outcomes	24
Panel Overview	25
Analysis	28
Event Study	28
Matched Controls	31
Synthetic Controls	34
Robustness	38
Interpretation	39
Conclusion	42
Bibliography	45
Appendix	47
A Primer on BAA and BABA	47
Exceptions to the Exception	48
Foreign Manufacture Share	49
Treated NSN Summary	50
Event Study Specification Ladder	51
Matched Controls: Per-NSN Distributions	56
Synthetic Controls: Per-NSN Detail	61
Robustness	64
Common Treated Sample	65
Excluding Freight Containers	69
Using Defense Logistics Agency Data Only	74

Acknowledgments

Two years ago, I had the incredible opportunity to intern at the Office of Management and Budget's Made in America Office. Being a part of this wave of industrial policy, and understanding the work that takes an idea to policy and then into reality, left a lasting impression on me. My experience inspired me to research the efficacy of public waivers as a tool of industrial policy, as part of this thesis. This work is my own, and does not represent the views of the Office of Management and Budget, the Made in America Office, the Biden Administration, or the Trump Administration.

I would like to express my gratitude to my advisor, Professor Conlon. Beyond his guidance at every stage, he broadened my research horizons, introducing me to approaches I had not encountered before. Professor Conlon was patient, generous with his time, and understanding throughout. Working with him made this past year a rewarding learning experience, and I am a better researcher for it.

I would also like to thank Professor White for his excellent leadership of the NYU Stern Honors Program, for the thoughtful feedback he offered on this thesis, and for the understanding and encouragement he showed throughout the year. Thanks to him, the Wednesday-evening seminars were a highlight of my senior year, an enriching experience that supercharged my interest in research.

I am deeply grateful to Eddie Garcia, and everyone else who took the time to talk with me, share their expertise, and provide their own insights on this subject. They helped me navigate through the weeds of federal policies, and their own interest motivated me all the more to pursue this thesis.

Thank you to my peers in the Honors Program, whose camaraderie softened the blow of whatever obstacles we ran into and made the good moments all the better for being shared. One day, we will finally understand Phased-Lock loop systems together.

Finally, I thank my family, who, despite being subjected to many long explanations of Buy American policies and F.A.R. clauses, provided their unwavering support to propel me past the finish line. Thank you.

Introduction

The federal government is a big buyer. In fiscal year 2025, federal agencies spent roughly \$793 billion on goods and services, making the United States government the single largest buyer in the world. For most of the past century, a Depression-era law called the Buy American Act has required this enormous purchasing power be directed toward domestic industry whenever possible.

But sometimes, items on the government’s shopping list are not manufactured in the United States. When this happens, the agency buying the product files a “nonavailability waiver”: a formal determination that no domestic supplier exists, allowing the purchase of a foreign-made product. For decades, every agency processed these exceptions internally, essentially burying the details within their procurement files, which made these exemptions practically invisible to outsiders.

In January 2021, Executive Order 14005 created the Made in America Office within the Office of Management and Budget. For the first time, nonavailability waivers from across the federal government were required to be submitted to a centralized portal and published for public review. The reasoning was that if the government publicly announces when it cannot source a product domestically, that information could help close the gap. Private industry might then notice the opportunity and step in to fill it. In addition, policymakers can use the information to improve outreach, market research, and other efforts to strengthen domestic supply. This process is “supply chain illumination”, where transparency about shortcomings in domestic production can itself serve as a tool of industrial policy.

Among the various instruments that have been deployed in recent years to strengthen domestic manufacturing, public waivers operate differently. Unlike tariffs, which raise costs for importers, or direct subsidies like those in the CHIPS and Science Act, which require large appropriations, the waiver portal’s mechanism is primarily informational and relatively low-cost. Public waivers convert what had been an internal administrative process into a public broadcast of unmet demand. The impact of the broadcast, and

whether it changes what gets procured and where, is the central question of this thesis.

To answer that question, I cross-reference two federal datasets: the Made in America waiver portal and data from USAspending.gov. The Made in America portal records every nonavailability waiver submitted since the system’s launch; USAspending.gov provides procurement contract data containing detailed records of federal purchases. The analysis examines Defense Logistics Agency supply chains, where product-level National Stock Numbers allow individual items to be tracked across time and across suppliers. By comparing procurement patterns for products that received nonavailability waivers against selected control products that did not, I evaluate whether the public disclosure of waivers has produced measurable shifts in domestic sourcing.

This analysis finds that the broadcast does not appear to be moving the supply chains in the intended direction. Across an event study, matched controls, and synthetic controls, domestic sourcing falls after a waiver is posted. The average number of offers received per transaction does not measurably change. Maximum unit price falls in the long term, though not immediately. The three designs handle pre-existing trends differently, but agree on the direction of each outcome.

Background

The History of Buy American

The Buy American Act was signed into law on March 3, 1933, by President Herbert Hoover the day before he left office. It emerged from the same wave of protectionist policy that had produced the Smoot-Hawley Tariff of 1930, as Congress sought to shield domestic manufacturers from foreign competition during the Depression. At its core, the Act required that when the federal government purchases goods, it should buy American-made products.

This basic principle has survived for over ninety years, though the details have been revised repeatedly. The Trade Agreements Act of 1979 carved out exceptions for products from countries with reciprocal trade agreements, creating a parallel track for national treatment, where certain foreign goods receive the same treatment as domestic ones. Over time, the Federal Acquisition Regulation (FAR, the sprawling rulebook that governs how the government buys things) codified the specific standards for domestic preference in its Part 25. These rules define what qualifies as a “domestic end product” (historically, a product with at least 55% domestic component content by cost, though that threshold has been tightened in recent years), how much of a price premium the government will pay for buying domestic (20% over foreign alternatives, or 30% for small businesses (“Buy American Act,” n.d.)), and under what circumstances exceptions are allowed.

Three broad exceptions permit agencies to bypass Buy American requirements: when domestic products are unreasonably expensive, when the purchase serves the public interest, and when the product is simply not available from domestic sources. This third category – nonavailability – generates the bulk of waivers and is the focus of this thesis.

All being said, this is still a simplified overview of a complex system. The full regulatory landscape of federal procurement has many more clauses, thresholds, and agency-specific supplements that are beyond the scope of this analysis.

Before Centralization

For most of the Buy American Act’s history, decisions around exceptions were typically handled by the contracting agency itself. Agencies worked within the government-wide FAR framework, but the steps could differ depending on the kind of exception, the size and structure of the purchase, and the agency involved.¹ In some cases a written nonavailability determination was required; in others it was not.²

A 2018 Government Accountability Office report (U.S. Government Accountability Office 2018) examined Buy American compliance across several federal agencies and illustrated how this system operated in practice. Agencies differed in the guidance and training they gave contracting officers, and GAO found that some contracting staff lacked complete guidance for required determinations and reviews. In the case of domestic nonavailability, the FAR set the basic rule, but agencies could still differ in who reviewed determinations and what documentation they expected.

For outside firms, that meant nonavailability decisions were not presented as standardized public cases. What became public was the procurement record after the fact: entries in the government’s FPDS procurement system showing that a foreign end product had been acquired and coded under a Buy American exception. GAO noted that only in 2018 were agencies’ historical reports on exception use added to fbo.gov. These reports essentially repackaged existing public information from FPDS, so vendors selling domestic products could more easily see which past contracts had fallen under Buy American exceptions.

This system showed which specific contracts relied on foreign-manufactured products. But those records were retrospective procurement entries, not specific explanations of the sourcing problem. They gave outside firms some signal that foreign supply had been used, but much less information about the unmet domestic need behind the purchase.

¹“The steps that contracting officers must take to determine or document an exception will vary depending on the circumstances of the acquisition.” (U.S. Government Accountability Office 2018).

²“For example, a written determination from the Head of Contracting Activity (HCA) or a delegate may be necessary to determine non-availability in some cases. However, a written determination may not be required when an acquisition is conducted through full and open competition, is synopsisized, and no domestic offer is received.” (U.S. Government Accountability Office 2018)

The New Waiver System

On January 25, 2021, President Biden signed Executive Order 14005, “Ensuring the Future Is Made in All of America by All of America’s Workers.” The order created the Made in America Office (MIAO) within the Office of Management and Budget and authorized it to review all proposed nonavailability waivers before they could take effect.³ Now, this part of the waiver process would be centralized.

The operational details were laid out several months later in OMB Memorandum M-21-26 (Young 2021). Under the new system, agencies proposing to waive Buy American requirements must submit their request to the MIAO, which posts the proposed waiver on [MadeinAmerica.gov](https://www.madeinamerica.gov). If the waiver passes review, the MIAO approves the waiver and the foreign purchase proceeds.

A foreign purchase could sometimes already be traced through FPDS and the underlying solicitation. But that information was scattered across procurement records and was not presented as a single, public nonavailability case. The new waiver system makes that determination explicit. Public waivers then reduce the cost of identifying and interpreting unmet government demand by linking the product, the agency’s stated rationale, duration of the waiver, and parts of the procurement context in one place.

On the Made in America website, illumination is an explicitly stated goal of the waivers: “This transparency portal provides access to information on proposed and approved waivers in a way that creates value for multiple stakeholders: Businesses can identify market opportunities where domestic manufacturing could fill existing gaps [;] Federal agencies can share information that signals demand for Made in America compliant products.” (“Waivers | Made in America,” n.d.)

The Biden administration also oversaw the rollout of a different type of public domestic procurement exemption, which is out of the scope for this thesis (see the Appendix entry on A Primer on BAA and BABA).

³The Made in America Office would later be codified into law with the passage of the Infrastructure Investment and Jobs Act. (Congress.gov. “Text - H.R.3684 - 117th Congress (2021-2022): Infrastructure Investment and Jobs Act.” November 15, 2021. <https://www.congress.gov/bill/117th-congress/house-bill/3684/text>.)

Literature Review

This thesis focuses on the impact of publicly broadcasting nonavailability waivers, whereas most analysis of Buy American regulation focuses on the impact of the regulation itself – the price preferences and rules directing federal spending to American manufacturers. The existing literature, however, provides important context for interpreting what the data shows.

The Cost of Buying American

The most directly relevant precedent is Bombardini, Gonzalez-Lira, Li, and Motta’s 2024 NBER working paper, “The Increasing Cost of Buying American.” Using procurement microdata from FPDS and a quantitative trade model, the authors estimate that existing Buy American provisions on final goods have created up to 100,000 domestic manufacturing jobs, at a cost of between \$111,500 and \$137,700 per job. The expense arises from higher procurement prices, market distortions, and the diversion of spending from more competitive foreign suppliers. This paper establishes that domestic preference policies can and do redirect purchasing toward American manufacturers, albeit at a fiscal cost.

Job creation, however, is one piece of the policy’s overall ambitions. Buy American steers procurement towards domestic production, and the case for doing so spans national security, supply chain resilience, and spurring innovation within American manufacturing ecosystems. With that motivation in mind, this thesis looks at how a new part of the Buy American regime – public waivers – fits into that effort.

Information Transparency and Sourcing Outcomes

Intuitively, greater transparency in the procurement process should lead to greater efficiency and better outcomes for buyers. Ohashi (2009), for example, studied Japanese municipal public works auctions and found that improved transparency in bidder qualification reduced procurement costs. If transparency helps identify where government demand is going unmet, then public waiver posting could improve the match between agencies seeking compliant products and domestic suppliers capable of producing them.

The Motivation Behind Public Waivers

The theory behind making all of this information public runs as follows: the newly public waiver process makes unmet domestic sourcing needs visible in a more usable form than before. By making waivers public, the government removes an informational barrier. The federal demand, now visible to anyone with an internet connection, serves as an invitation for firms to invest in production capacity and compete for future contracts, and highlights industry gaps to procurement officials.

This is what policymakers mean by “supply chain illumination.” The waiver portal itself does not subsidize production, does not restrict imports, and does not mandate entry. Instead, it makes a specific kind of information freely available: *the government needs this, and domestic supply is insufficient.*

Imagine an American company is considering entering the federal procurement market with a certain domestically-manufactured product. The company does its due diligence, and sees that the federal government is already buying that product from a foreign manufacturer, citing domestic nonavailability.

Under the old system, that’s the end of the story. The company has no information beyond the nonavailability citation – it does not know if the government’s foreign purchase was a one-time deviation, if the government will use the same exception the next time it buys the product, or if the agency is obligated to conduct full market research for its next purchase which may lead it to the domestic company. Unsure, the company decides to move on, and the American-made product goes unsold.

Things may go differently under the new system. The American company does its due diligence and sees the purchase had received a single-use waiver, only applicable for that specific purchase. It knows the next time the agency wants to buy the product, the waiver will no longer apply, meaning the company will have a better chance of landing a contract. The company decides the investment in becoming a government supplier is worth it, and now the agency has a domestically-manufactured option available.

Illumination relies on some conditions to work.

First, the right people must see the data. Federal contractors routinely monitor SAM.gov (System for Award Management) for procurement opportunities. MadeinAmerica.gov is a separate website with a different purpose and a smaller audience. If manufacturers are not actively checking the waiver portal, the demand signal may not reach its intended recipients.

Second, federal demand needs to be large enough to change sourcing behavior. The government is an enormous buyer in aggregate, but its purchasing power varies dramatically by product category. For specialized defense items where the government is effectively the only customer, a posted waiver could be decisive. For commodity products like medical supplies⁴, where government purchases represent a small fraction of total market demand, the signal may be too weak to change a firm's investment calculus.

⁴Combined, the Department of Defense and Department of Veteran Affairs account for 2 to 3 percent of US medical-supply demand, according to congressional testimony from then-Executive Director for Procurement and Logistics at the Veterans Health Administration, Andrew Centineo. (Centineo testimony, VA Procurement: Made in America, September 20, 2023)

Waiver Hypotheses

The central hypothesis is that public waiver disclosure changes subsequent procurement outcomes for the affected product. This can be broken into three predictions around domestic sourcing, competition, and price.

Hypothesis 1: Domestic Sourcing

Public waiver disclosure may shift the sourcing for the affected product toward domestic manufacture. If domestic suppliers enter or expand their participation in response to the demand signal, the share of transactions manufactured in the United States should rise relative to the pre-waiver baseline. This follows the overall waiver policy’s stated purpose.

Hypothesis 2: Competition

Public posting of a nonavailability waiver may increase the number of offers received for that product in subsequent procurements. By revealing unmet demand to potential domestic suppliers, the waiver draws new bidders into the procurement process. Even the partial entry of firms that bid but don’t win would register as an increase in offers received.

Hypothesis 3: Price

Bulow and Klemperer (1996) show that adding a bidder to a competitive auction compresses the price distribution and benefits the buyer. That means more entry ought to lower prices. But DoD procurement complicates this prediction. Under DFARS (the Defense Department supplement to FAR), contracting officers apply a 50% evaluation factor to foreign bids when comparing them against domestic offers⁵⁶.

For example, a foreign bid of \$1,000 is treated as \$1,500 for award purposes. A domestic supplier drawn in by a posted waiver can therefore win a contract while pricing up to

⁵DFARS 225.106(b): “Use an evaluation factor of 50 percent instead of the factors specified in FAR 25.106(b)(1)(i) and (c)(1)(i).” The factor is applied to the low foreign offer per DFARS 225.502(c). (“Subpart 225.1 - BUY AMERICAN—SUPPLIES | Acquisition.GOV,” n.d.)

⁶This is notably more than the 20% to 30% price preference offered by civilian agencies (“Buy American Act,” n.d.)).

50% above the foreign alternative. Rather than driving prices down, domestic entry could raise what the government pays, achieving the sourcing goal at a fiscal cost. This trade-off would be consistent with the broader findings in Bombardini et al. (2024). As a result, public waivers may raise the unit costs of products for agencies.

Research Design

The treatment is the public posting of a nonavailability waiver on MadeinAmerica.gov. Federal agencies have bought foreign-made products under Buy American exceptions for decades; 2021’s Executive Order 14005 required these written determinations to be posted publicly and centrally⁷. The moment the waiver appears on the public portal, the government’s unmet domestic demand for that product becomes visible to anyone checking the site.

Unit of Analysis

The analysis is conducted at the level of the National Stock Number (NSN), a 13-digit identifier used by the Department of Defense and the Defense Logistics Agency to catalog individual supply items. Each NSN identifies a specific product and remains constant over time regardless of who manufactures it or which office purchases it. NSN 4240-01-538-7970, for example, covers a specific type of hearing protector.

NSNs make product-level tracking possible in a way that broader classification schemes cannot. Product Service Codes (PSC) are widely available in government procurement data, but are too broad to discern specific products. For example, PSC 4240 is titled “Fire/Rescue/Safety/Environmental Protection Equipment” and encompasses hearing protectors, gas masks, and breathing apparatus valves.

Because an NSN refers to a single cataloged item, procurement outcomes can be traced for the exact product named in the waiver across a multi-year window. This restricts the analysis to DoD and DLA waivers, since civilian agencies do not use the NSN system.

⁷Like the establishment of the Made in America Office, this requirement for written determinations was then codified into law by the Infrastructure Investment and Jobs Act. (Congress.gov. “Text - H.R.3684 - 117th Congress (2021-2022): Infrastructure Investment and Jobs Act.” November 15, 2021. <https://www.congress.gov/bill/117th-congress/house-bill/3684/text>.)

Treatment Panel

The treatment panel consists of 27 NSNs with both pre-waiver and post-waiver procurement records, including at least one observation in the five-year window before the waiver posting date (event years -5 through -1). Twenty-one are DLA freight container NSNs. The remaining six are individual items in various product categories: a heat exchanger, an airframe warning system component, a pneumatic motor sourced under French NATO cataloging, a set of material handling rollers, a hearing protector, and a shower bath fixture. The treatment event for each NSN is the date its waiver first appeared on [MadeinAmerica.gov](https://www.madeinamerica.gov).

The pre-treatment period extends back to fiscal year 2017. Waiver posting dates fall within 4 fiscal years (FY2022 to FY2025), and the staggered timing helps distinguish waiver effects from common time-variant shocks. There are three outcomes measured at the NSN-by-event-year level:

- Domestic sourcing share: the share of transactions where the product was manufactured in the US.
- Number of offers received: the average number of bids per transaction.
- Unit cost: the maximum log unit price.

The thesis uses three estimators. The first is an event study with NSN and fiscal-year fixed effects. The second is a matched difference-in-differences, which pairs each treated NSN with similar untreated NSNs before comparing pre- and post-waiver outcomes. The third is synthetic difference-in-differences (Arkhangelsky et al. 2021), which builds a custom synthetic comparison for each treated NSN from a weighted combination of untreated NSNs.

Beyond accounting for pre-existing trends, both matched and synthetic designs compare treated NSNs to products with similar procurement histories, instead of the full catalog of Defense procurement. Both also report effects at the level of the individual NSN, as well as in aggregate. Since freight containers make up most of the treatment panel, findings

from these estimators can help reveal whether a pooled result is observed across treated NSNs, or driven by a handful of influential items.

All three estimators use the same panel structure. Each observation is one NSN in one fiscal year. Fiscal years that straddle a treated NSN’s waiver anniversary are split in two, so that the pre- and post-waiver portions of that fiscal year can be analyzed separately. The Data section describes how the panel is built.

Event Study

The event study models each panel observation’s outcome as a function of NSN fixed effects, fiscal-year fixed effects, and event-year dummies.

$$y_{it} = \alpha_i + \gamma_t + \sum_{k \neq -1} \beta_k \mathbb{1}\{e_{it} = k\} + \varepsilon_{it}$$

In this specification, the outcome y is observed for each NSN i in each fiscal year t . The NSN fixed effect α_i absorbs differences between NSNs that do not change over the study period, like intrinsically higher price levels for certain products. The fiscal-year fixed effect γ_t absorbs shocks that hit all NSNs in a given year, such as a continuing resolution that delays procurement spending across the government.

Each event-year indicator equals one when NSN i is at event year k in fiscal year t , and the coefficient β_k measures the outcome at event year k relative to event year -1 , which is omitted as the reference. The error term ε_{it} captures variation in the outcome that the fixed effects and event-year dummies do not explain, and standard errors are clustered at the NSN level. The output is an event-time path: one coefficient per event year, with $k = -1$ normalized to zero.

Matched Controls

To address the question of pre-treatment trends driving the results seen in the event study, matched controls build a comparison group whose pre-waiver behavior already resembles

the treated NSNs'. For each treated NSN, the three most similar untreated NSNs in the pre-waiver period are selected, and the treated NSN's post-waiver change is compared to a weighted average of the changes among those three matches. Similarity is measured by Mahalanobis distance, a multivariate distance that puts variables on a common scale by accounting for their covariance structure. The matching variables are four summary statistics from the pre-waiver period: log transactions per fiscal year, average maximum log unit price, average domestic share, and average number of offers received. Matching is with replacement, so a single untreated NSN can be matched to multiple treated NSNs.

For each treated NSN, the per-NSN treatment effect is the treated NSN's pre-to-post change minus the weighted average of the matched NSNs' pre-to-post changes:

$$\hat{\tau}_i = (\bar{y}_i^{\text{post}} - \bar{y}_i^{\text{pre}}) - \sum_j w_{ij} (\bar{y}_{ij}^{\text{post}} - \bar{y}_{ij}^{\text{pre}})$$

Here, τ_i is the per-NSN treatment effect for treated NSN i . The terms $\bar{y}_i$ -pre and $\bar{y}_i$ -post are the average outcomes for the treated NSN before and after the waiver. $\bar{y}_{ij}$ -pre and $\bar{y}_{ij}$ -post are the analogous pre- and post-waiver averages for matched untreated NSN j . The weights w_{ij} are produced by the matching step, which selects the three nearest neighbors by Mahalanobis distance over the four pre-period covariates described above. The pooled average treatment effect on the treated reported in the body is the simple mean of τ_i across treated NSNs.

Synthetic Controls

Synthetic difference-in-differences (Arkhangelsky et al. 2021) extends the synthetic control method by combining unit weights with time weights inside a two-way fixed-effects regression. For each treated NSN, the estimator solves the following weighted least-squares problem:

$$(\hat{\tau}, \hat{\mu}, \hat{\alpha}, \hat{\beta}) = \arg \min_{\tau, \mu, \alpha, \beta} \sum_{i=1}^N \sum_{t=1}^T (Y_{it} - \mu - \alpha_i - \beta_t - W_{it} \tau)^2 \hat{\omega}_i \hat{\lambda}_t$$

In this expression, Y_{it} is the outcome for NSN i in fiscal year t , and W_{it} equals one when NSN i is the treated NSN and t falls in the post-waiver window. The terms α_i and β_t are NSN and fiscal-year fixed effects, and μ is a common intercept. The unit weights ω_i are zero for the treated NSN and non-negative across the donor pool, chosen so that the weighted average of donor pre-waiver trajectories tracks the treated NSN's pre-waiver trajectory. The time weights λ_t are zero in post-waiver periods and chosen so that the weighted average of pre-waiver outcomes approximates each donor's post-waiver mean.

The estimated treatment effect τ is therefore the post-waiver gap between the treated NSN and its synthetic comparison, net of any pre-waiver level difference absorbed by the time weights. Eligible untreated NSNs are restricted to those that observe the outcome in every event year the treated NSN does, which prevents the synthetic from being constructed by extrapolating across years of missing data.

Data

The previous section described the research design in broad strokes: track procurement outcomes for products that received public nonavailability waivers, compare them to similar products that did not, and measure whether public disclosure changed procurement outcomes. This section describes the data that make that design operational, including what the sources are, what the broader waiver and procurement landscape looks like, how the analysis sample gets assembled from it, and how outcomes are constructed from the joined records.⁸

Sources

The analysis draws on four data sources.

The Made in America Office publishes waiver submissions on [MadeinAmerica.gov](https://www.madeinamerica.gov). Each waiver record includes the contracting and funding agency, a text summary of the procurement, the Product Service Code, the waiver’s expected duration, and its review status. The dataset used here contains over 2,000 waivers observed through March 2026. Nearly all are nonavailability determinations; the remaining two are public interest waivers. The earliest records date to November 2021, when the portal launched.

Contract-level procurement records come from [USAspending.gov](https://www.usaspending.gov), which publishes data aggregated from various government databases ([USAspending.gov](https://www.usaspending.gov) 2023, 8–9). The downloaded data covers fiscal year 2017 through January 2026 and contains approximately 37.7 million transaction records after excluding non-manufactured end products.⁹

USAspending records each modification under a contract, meaning each row has a unique contract transaction key, and rows under the same contract share the same Procurement Instrument Identifier (PIID). For the analysis of NSN transactions, only the original contract transaction is used; future modifications are ignored (see Building the Analysis Panel

⁸Code used for this thesis is available at <https://github.com/Anantesh-Mohapatra/Waived-in-America>.

⁹Transactions coded with place of manufacture code “C” (items not manufactured as end products, such as services coded under a manufacturing PSC) are excluded from the procurement extract. These contracts are not subject to the Buy American Act.

section). Each record includes the awarding and funding agency, the contractor and its CAGE code, the place of manufacture, the country of product origin, the number of offers received, the dollar obligation, and a text description of the purchase.

There are two limitations with this dataset. First, USAspending does not include a dedicated National Stock Number field, since NSNs are more relevant for military procurements. Second, it does not record the quantity purchased per transaction, only dollar obligations at the contract-modification level, so unit prices cannot be derived from USAspending alone.

The Defense Logistics Agency publishes its contract history through its FOIA reading room. Each record is a line item from a DLA-administered contract, with the NSN, contract identifier, award date, contractor CAGE, order quantity, and unit price. The DLA data used in this thesis covers fiscal year 2017 through January 2026 and contains approximately 13.4 million line-item records. BidLink, a commercial procurement database, extends coverage for waived items and a subset of non-treated NSNs.

Descriptive Statistics for Waivers

The over 2,000 nonavailability waivers posted since the portal’s launch cover everything from Phase-Locked Loop systems to freight containers. Civilian science and health agencies dominate. The National Institute of Standards and Technology leads with 517 waivers, followed by the Department of Veterans Affairs (418), the National Institutes of Health (394), NOAA (114), and the Agricultural Research Service (105), as of March 9, 2026 (when the waiver dataset was downloaded). Together with the CDC, these agencies account for more than three-quarters of all waivers posted.

Most waivers are short-lived. About 58 percent are classified as “instant delivery only,” covering a single purchase already in progress, and another 21 percent cover six months or less. Multi-year waivers are uncommon, at about 7 percent of the total. The predominance of instant-delivery waivers means that many records formalize an individual transaction rather than authorizing a sustained program of foreign purchasing.




Figure F1. Nonavailability waivers granted by contracting agency and fiscal year.

This stems from a regulatory preference to have more short-term waivers to provide clearer insights into federal procurement patterns. Hypothetically, long open-ended waivers can dilute the demand signal (since they aren't tied to specific procurements) and reduce the incentive for agencies to procure domestically manufactured products (since future procurements wouldn't require the friction of requesting a waiver).

Product Service Code 6640, laboratory equipment and supplies, covers nearly half of all waiver records. Medical products, drugs, and scientific instruments make up most of the remainder. At the product level, the most granular information provided by these waivers are Product Service Codes (PSC) and product names included in the metadata. In the absence of a specific catalog number or product identifier, product-level analysis of these waivers is difficult.

The Department of Defense accounts for about 205 waivers, or roughly 10 percent of the total. Within DoD, the Defense Logistics Agency filed the most (81), followed by the Navy (62) and the Army (54). From this subset, only waivers disclosing explicit National Stock




Figure F2. Nonavailability waivers by expected duration and fiscal year.




Figure F3. Nonavailability waivers by product/service code and fiscal year.

Numbers can be tracked at the product level in the procurement data. There are 44 such waivers.

Procurement in Context

To put the waivers in perspective, consider the broader landscape of federal procurement.

Of the roughly 37.7 million federal procurement transactions for manufactured end products recorded between FY2017 and January 2026, the vast majority (97 percent) involved domestically manufactured products. Foreign-sourced transactions account for only 2 to 3 percent of all transactions in any given year, though they represent a larger share of total dollars spent, ranging from about 3 percent in FY2018 to nearly 6 percent in FY2022.

Nonavailability determinations (place of manufacture code “J” in FPDS) represent about 0.2 to 0.5 percent of all transactions annually. Within the foreign-manufacture category, however, nonavailability is far more prominent: about 6 percent of foreign transactions in FY2017-2018, rising to between 11 and 23 percent from FY2019 onward (Figure A0).

Nonavailability-coded procurement climbed to roughly \$4.2 billion in FY2022 (about 1.7 percent of all contract dollars) before moderating in subsequent years. This increase was driven by purchases of COVID test kits from China and South Korea, and COVID vaccines from Denmark and Canada.

The two other major categories of foreign procurement are Trade Agreements Act purchases (code “G”) and purchases for use outside the United States (code “E”). These do not require individual nonavailability waivers and are not the focus of this analysis.

Within this picture, DLA accounts for the bulk of federal procurement of manufactured end products recorded in FPDS – roughly 25 million transactions over the period, which is about two-thirds of all such transactions. DLA’s domestic sourcing share is higher than the federal average (98 percent of transactions versus 97 percent), and its nonavailability rate is correspondingly lower (0.09 percent of DLA transactions, compared with the 0.35 percent federal-wide rate).

This means domestic sourcing is the overwhelming norm in federal procurement. Foreign sourcing is the exception, and nonavailability sourcing is a subset within it. The dollar value of nonavailability-coded procurement amounts to billions of dollars per year. This spending category is targeted by the waiver system.

From Waiver Universe to Analysis Sample

The analysis sample is much smaller than the full waiver universe. This is because there are two major problems with the two main datasets of this thesis.

The first issue is that the waiver data and the procurement data do not always share a common key. There is no field in the standard USAspending schema that reliably links a specific procurement transaction to a specific waiver. When a procurement is nonavailability-coded, it does not have a corresponding waiver key or other nonavailability justification (see Appendix ‘Exceptions to the Exception’). And vice-versa, most waivers are not linked to a specific PIID that would allow viewers to find corresponding contracts. Because waivers do not publicly disclose the cost of the product being waived, a PIID is needed to access additional details surrounding the procured product.

The second issue is around product identification. Both waivers and USAspending procurement records reliably have PSC codes – information on the broad product category. Both also have free-text descriptions of the product or procurement. But these descriptions are too inconsistent across waivers and contracts to support reliable product matching.

National Stock Numbers help serve as a bridge. Waivers sometimes contain explicit NSNs, and the DLA dataset can return individual contracts for a given NSN. With NSNs, the match between the waived product and procured item becomes exact. But this match is rare. Of the more than 2,000 waivers in the portal, 44 contain explicit NSN codes. All 44 are from the DLA or Navy¹⁰. The roughly 90 percent of waivers filed by civilian agencies like NIST, VA, and NIH cannot be tracked at the product level through NSNs, because civilian agencies do not use the NSN system.

¹⁰Some waiver records cover multiple NSNs while some NSNs appear in multiple waiver records filed at different times. For example: there is a freight container waiver that lists 20 container NSNs.

Table 1. Sample Funnel: From Waiver Universe to Treatment Panel

Stage	Count	Unit	% retained
All cleaned waivers	2,134	waivers	—
Nonavailability waivers only	2,132	waivers	99.9%
DoD / defense agency	204	waivers	9.6%
Waivers with explicit NSN codes	44	waivers	21.6%
Unique waiver NSNs	37	NSNs	84.1%
NSNs with any transaction history	31	NSNs	83.8%
With transactions both pre- and post-waiver	28	NSNs	90.3%
With pre-waiver transactions inside the analysis window (event years -5 to -1)	27	NSNs	96.4%

Stages 1-4 count waivers; stages 5-8 count NSNs. Stage 7 requires transactions on both sides of the waiver date so a pre/post comparison is possible. Stage 8 requires that pre-waiver transactions must fall within the 5-year analysis window.

Those 44 waiver records yield 37 unique NSN-format identifiers. Of these, six have no procurement records in either BidLink or DLA contract history, leaving 31 NSNs with at least some procurement history.

Not all 31 NSNs can support a before-and-after comparison. To qualify as a treatment unit, an NSN must have at least one transaction record in BidLink or DLA contract history before its waiver posting date and at least one after. Three NSNs fail this test. All three have pre-waiver procurement histories but no post-waiver transactions, and all three are post-solicitation waivers: the agency was already purchasing the item before the waiver was publicly disclosed. An additional waived NSN is eliminated for having no transactions in the five years before its waiver was posted.

These filters funnel over 2,000 waivers to 27 usable NSNs. Most federal nonavailability waivers describe products using commercial identifiers that cannot be followed through DoD procurement systems. Even among DoD waivers that do carry NSN codes, some lack the procurement history needed for a before-and-after comparison. The 27 treatment NSNs are the subset of defense supply items for which product-level procurement tracking is feasible.

Building the Analysis Panel

The analysis panel covers the 27 treated NSNs and a broader set of non-treated NSNs used for comparison. The DLA contract history, supplemented by additional transactions captured in BidLink, is joined to USAspending on the contract identifier to attach FPDS fields for competition, offers, place of manufacture, and obligation.

Each transaction has a fiscal year (its calendar position relative to the federal fiscal-year calendar) and an event year (its position relative to the treated NSN’s waiver date). Cells are the intersection of these two dimensions for a given NSN, so each cell holds the transactions that share the same NSN, the same fiscal year, and the same event year.

Consider an NSN waived in March 2023. Its event year 0 spans March 2023 through March 2024 and crosses the federal fiscal-year boundary on October 1, 2023. Transactions in event year 0 are therefore split into two cells: one encompassing the March 2023 to September 2023 transactions (which fall in fiscal year 2023), and another holding the transactions between October 2023 to March 2024 (which fall in fiscal year 2024).

Measuring Outcomes

Domestic sourcing share comes from USAspending’s place-of-manufacture field, which assigns each transaction one of twelve codes: D for US-manufactured products, and others distinguishing foreign sources by their regulatory basis (nonavailability, Trade Agreements, qualifying country, etc.). For each event year, the share is the number of D-coded transactions over the number with any manufacture code recorded.

The number of offers received comes from USAspending’s number-of-offers-received field. The event-year mean averages over transactions with a recorded offer count.

Unit prices come from DLA’s contract line items; the event-year outcome is the maximum log unit price across the event year’s transactions. For two NSNs, no direct unit price is recorded. For one of them – the heat exchanger – a proxy price is recovered from the contract obligation divided by total quantity ordered. This method is used only when a

contract covers a single NSN. The other, the airframe component, has too few single-NSN contracts and is excluded from the unit-price analysis.

Panel Overview

The full DLA panel covers about 947,000 NSNs over 9.0 million transactions in fiscal years 2017 through January 2026. For an NSN to contribute to the comparison, it must have at least two observations, including one in fiscal year 2022 or later (when waivers began to be issued). 372,000 NSNs meet this requirement, covering 7.6 million transactions – roughly 40 percent of the universe by NSN count and 85 percent by transaction count. Domestic sourcing share across both groups sits steadily near 97 percent over the study window, and a pool restricted to the same FSCs as the treated NSNs sits near 98 percent (Figure F4a). The filters and the FSC-level restriction do not meaningfully change the DLA-wide picture.




Figure F4a. Domestic sourcing share by fiscal year, treatment panel versus comparison groups.

The 27 treated NSNs are a much smaller slice of DLA transactions: roughly 2,200 transactions across 21 freight containers and six individual items. Per-NSN counts range from

a handful of transactions to several hundred (Table A1).

The treated panel’s fiscal-year trajectory diverges sharply from every comparison line in Figure F4a. Domestic share runs near full through fiscal year 2019, falls into the 40-70 percent range during 2020-2022, and reaches single digits by 2023. The same-FSC pool, by contrast, stays near 98 percent throughout.

By volume, the treated NSNs shift drastically to Chinese manufacturers. Chinese manufacturing overtakes domestic sources in fiscal year 2020 (before the introduction of the public waiver policy), although American and South Korean manufacturers retain a large share of spending until fiscal year 2025 (Figure F4b). Figure F4c shows that this shift isn’t driven only by the containers, and occurs in the non-container products as well.




Figure F4b. Country of origin across treated NSNs by fiscal year.

Looking at the average domestic sourcing trajectory of treatment NSNs, relative to their waiver posting date (Figure F5), the decline begins around event year -5 and reaches roughly 5 percent at event year 0, then recovers to about 10 percent in the following years. To address this pre-treatment trend, the matched-controls and synthetic-control estimators match on pre-waiver behavior.




Figure F4c. Country of origin over time for the six non-container treated NSNs.




Figure F5. Domestic sourcing share, treatment panel averaged by event year.

Analysis

Event Study

All 27 treated NSNs enter the event study, alongside over 300,000 non-waived NSNs from the filtered DLA universe that have at least one cell in FY 2022 or later (when the public waiver policy came into effect) and at least two recorded observations per outcome. As described in the Research Design section, the regression has NSN and FY (fiscal year) fixed effects, event-year dummies with event year -1 as the reference, and standard errors clustered on NSN.

Specifications adding fixed effects one at a time appear in Appendix Tables A2-A4 and Figures A1-A3.

Domestic sourcing share shows large negative coefficients at every post-waiver event year: -0.285 (SE 0.073) at event year 0, -0.213 (SE 0.068) at EY+1, -0.237 (SE 0.089) at EY+2, and -0.314 (SE 0.065) at EY+3. These are percentage-point changes in the share of an NSN's transactions where the product was manufactured in the US, so that share is 21 to 31 percentage points lower at each post-waiver event year than at event year -1, after absorbing NSN and FY effects. Each rejects zero at $p < 0.05$ individually, and the linear combination across the post-waiver event years is -0.262 (SE 0.065, $p < 0.01$).

The pre-treatment coefficients, however, are also large. Both visually (Figure F6) and by the joint F-test for the pre-treatment dummies being zero ($p < 0.01$), the assumption of parallel trends is rejected. These products were already trending toward non-domestic sourcing before their waivers were posted, so it's harder to attribute the post-waiver change to the waiver itself. The matched-controls and synthetic-controls designs that follow both build their comparison groups by matching on pre-waiver behavior, which directly addresses this.

Maximum logged unit price shows a cleaner story. The pre-waiver trend is flat, with the joint F-test for the pre-treatment dummies returning $p = 0.462$. The coefficient for event year 0 is not significant (-0.077, SE 0.082), but the following post-waiver coefficients are




Figure F6. Event-study coefficients for domestic sourcing share.

uniformly negative and significant ($p < 0.01$): -0.150 (SE 0.053) at EY+1, -0.187 (SE 0.063) at EY+2, and -0.302 (SE 0.116) at EY+3 (Figure F7).

The outcome is the natural log of the maximum unit price observed for an NSN in each event year, and each coefficient is the change relative to event year -1. So -0.179 (the post-period average) corresponds to roughly a 16% lower maximum unit price than at event year -1 ($e^{(-0.179)} - 1 \approx -0.164$), and -0.302 (EY+3) to roughly a 26% lower maximum price. As with the individual event-year coefficients, the linear combination across post-waiver event years clears the 0.05 threshold for significance (-0.179, SE 0.062). In other words, the event study finds a sustained, statistically significant reduction in the maximum logged unit price.

Mean number of offers received produces post-waiver coefficients of -0.791 (SE 0.703) at event year 0, -0.688 (SE 0.394) at EY+1, -0.534 (SE 0.323) at EY+2, and -0.237 (SE 0.805) at EY+3 (Figure F8). These are changes in the average number of offers per transaction for an NSN in each event year, so the event year 0 estimate of -0.791 means about 0.8 fewer




Figure F7. Event-study coefficients for maximum logged unit price.

offers per transaction than at event year -1. None reach the 0.05 threshold individually, and the linear combination across the post-waiver event years is -0.562 (SE 0.432, $p = 0.19$), also not statistically distinguishable from zero.

The pre-waiver coefficients are also large and reject parallel trends. The offers regression runs on a smaller sample than the other two outcomes (about 1.12 million cells against 1.47 million) due to the availability of offers submitted in FPDS data.¹¹

The event-study results come with certain caveats. First, the pre-trend on domestic share poses an obstacle to interpreting the waiver effect. The NSN fixed effect removes time-invariant differences in level between treated and non-treated NSNs, and the FY fixed effect removes shocks that hit every NSN equally in a given year. Neither absorbs a trend that is specific to the treated NSNs and was already in motion before the waiver. So the post-waiver coefficients on domestic share reflect both whatever effect the waiver had and

¹¹Many code-A (full and open competition) awards are reported in FPDS without a populated number_of_offers_received field. About 30% of (NSN, FY, event year) cells in the non-treated universe have missing mean offers fields, versus 1.7% of treated cells.




Figure F8. Event-study coefficients for mean offers.

the continued descent of that pre-existing trend, and the event-study specification cannot distinguish between the two.

Second, not every waiver has been posted long enough for the later event years to be observed in the data, so the later event years are identified by fewer treated NSNs than the earlier ones. The confidence intervals at event year +2 and especially +3 widen accordingly in the coefficient plots.

Matched Controls

The matched-controls design fits 26 of the 27 treated NSNs. The one drop is 1680-01-622-9189 (an aircraft warning system component), which has no recorded pre-waiver unit price and so cannot supply the mean maximum logged unit price matching covariate. The eligible non-treated pool starts from the same 372,288-NSN filtered DLA universe used by the event study, then narrows further by requiring that any candidate match observe the outcome in every event year the treated NSN does – the same balanced-coverage rule that

Research Design describes for the synthetic-controls design. Per-treated pool sizes after this rule range from roughly 8,800 to 82,000 non-treated NSNs.

For each treated NSN, the matched-controls design produces a per-NSN treatment effect τ_i : the change in that NSN’s outcome from pre-waiver to post-waiver, minus the weighted average of the same change among its three matched non-treated NSNs. The average treatment effect on the treated (ATT) is the mean of the 26 τ_i values, with HC1 (heteroskedasticity-robust) standard errors.

Two post-period windows are reported. The first, event year ≥ 0 , uses all 26 treated NSNs. The second, event year ≥ 2 , keeps only the 14 treated NSNs whose post-waiver records extend at least two years beyond their waiver date (13 for mean offers¹²). Both windows are also estimated under a same-FSG sensitivity, which requires each donor to share the treated NSN’s two-digit Federal Supply Group (the first two digits of the four-digit FSC code) – enforcing a degree of product similarity. All 26 treated NSNs still match under this constraint; Table 2 reports both variants side by side.

Table 2. Matched Controls: ATT by Outcome and Post-Period Window, Main Matching vs. Same-FSG Sensitivity

Outcome	Post window	N	Main matching ATT (SE)	Same-FSG sensitivity ATT (SE)
Domestic sourcing share	event year ≥ 0	26	-0.397*** (0.071)	-0.299*** (0.070)
	event year ≥ 2	14	-0.226* (0.111)	-0.271** (0.101)
Maximum logged unit price	event year ≥ 0	26	-0.117 (0.086)	-0.126 (0.095)
	event year ≥ 2	14	-0.170** (0.058)	-0.224*** (0.045)
Mean offers	event year ≥ 0	26	-0.494 (0.719)	-0.086 (0.680)
	event year ≥ 2	13	-0.822 (0.486)	-0.007 (0.317)

ATT is the mean of the per-NSN τ_i , with HC1 (heteroskedasticity-robust) standard errors. *Main matching* picks the three nearest donors by Mahalanobis distance from the full eligible pool. *Same-FSG sensitivity* adds an exact-match constraint requiring each donor to share the treated NSN’s two-digit Federal Supply Group (the first two digits of the four-digit FSC code); all 26 treated NSNs still match under this rule.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

¹²One of the 14 treated NSNs is dropped, as all three of its matched donors have blank offer counts on their event-year ≥ 2 transactions.

On domestic sourcing share, the ATT is -0.397 ($p < 0.01$) under the full window and -0.226 ($p < 0.10$) under $EY \geq 2$ (Table 2). The -0.397 ATT means that, comparing pre-waiver to post-waiver, the average treated NSN’s share of transactions manufactured in the US drops roughly 40 percentage points more than its matched donors’ share does. As Figure A4 shows, the per-NSN distribution is negative for nearly every fit; only two τ_i are positive, and one is exactly zero. The largest negative τ_i is -1.000 for three NSNs whose domestic share fell from one to zero across the window – a freight container, the hearing protector, and the pneumatic motor.

Maximum logged unit price differs between the two windows. Under the full post window, the ATT is negative but does not clear the 0.05 threshold. Under $EY \geq 2$, the ATT is -0.170 ($p < 0.05$) – on the natural-log scale, that means the average treated NSN’s maximum unit price drops about 16% more from pre-waiver to post-waiver than its matched donors’ does. The per-NSN distribution sits mostly below zero with one large positive outlier – the hearing protector NSN at +1.77, the same NSN that drives the largest domestic-share drop (Figure A6).

The mean of offers produces ATTs that are not statistically distinguishable from zero in either window, with SEs below one offer-per-transaction. The per-NSN τ_i span -13.2 to +10.0 (Figure A8), driven by a handful of NSNs whose pre-waiver offer counts were already in the double digits.

The matched-controls estimator has post-match imbalances on two covariates. Standardized mean differences sit below the 0.25 threshold for log transactions per fiscal year (mean $|SMD|$ 0.14) and mean number of offers received (0.16). Mean domestic share is at the threshold (0.25), with 7 of the 26 matched pairs above 0.25 individually, and maximum log unit price exceeds it (0.28), with 12 of 26 pairs above.¹³

¹³The standardized mean difference compares the treated and donor means relative to their pooled standard deviation.

Synthetic Controls

The synthetic-controls design fits 22 treated NSNs on maximum logged unit price, 22 on mean offers, and 23 on domestic share. Of the 27-NSN starting set, synthdid identification requires at least three pre-waiver event years for each treated NSN; four NSNs whose pre-waiver procurement history is too thin fall below this threshold and are excluded. Synth’s unit-price fits also exclude the aircraft warning system component that matched-controls dropped, for the same reason: no recorded pre-waiver unit price.

The synthdid package (Arkhangelsky et al. 2021) is used for estimation. Each (treated NSN, outcome) is fit separately against its own non-treated pool. The donor-side rule is the same as in matched controls – each donor must have observations in every event year the treated NSN has observations in for that outcome – but because synth’s treated NSNs have already been filtered to have at least three pre-waiver event years, donors here must cover more pre-waiver periods than they did in matched. Pool sizes vary by outcome: median pool size is around 14,000 to 16,000 for unit price and domestic share, and around 3,000 for mean offers (FPDS often leaves the offer count blank, as noted in the event-study section, so fewer donors qualify under the coverage rule).

Inference uses synthdid’s placebo procedure. For each (treated NSN, outcome) estimate, 200 placebo runs are constructed: each randomly picks a non-treated NSN, treats it as if it were the treated one, refits the synthetic control, and records the resulting placebo ATT. The standard deviation of the 200 placebo ATTs gives the standard error, and the share of placebos whose absolute ATT meets or exceeds the observed ATT gives the empirical p-value.

Table 3. Synthetic-controls fit diagnostics by outcome.

Outcome	N fits	Median pre-fit RMSPE	Median active donors / pool size
Domestic sourcing share	23	0.0032	82 / 14,024
Maximum logged unit price	22	0.0004	4,730 / 15,568
Mean offers	22	0.0076	1,877 / 3,057

ATTs are estimated per treated NSN using `synthdid` (Arkhangelsky et al. 2021). Significance is assessed against the random-subsample placebo distribution (200 placebos per fit). Pre-fit RMSPE is the root mean squared prediction error in the pre-waiver period — smaller means the synthetic tracks its treated NSN more closely before the waiver. The “active donors / pool size” column reports the median count of non-treated NSNs that receive non-trivial weight ($w > 1 / (2 \cdot \text{pool size})$) against the total eligible pool — lower ratios indicate that the synthetic relies on a small subset of the pool.

On domestic sourcing share, the per-NSN ATTs lie on the negative side, and 21 of the 23 per-NSN estimates are individually significant at $p < 0.05$ (Table A5, Figure A10). Two estimates sit close to zero and no NSN has a significantly positive estimate. The -0.46 mean at event year 0 means that, on average, treated NSNs have a share of transactions manufactured in the US that is 46 percentage points below their synthetic comparisons at event year 0.

As Figure A11 shows, maximum logged unit price gives a mixed per-NSN picture: the median estimate is negative, the cross-NSN standard deviation is 0.44 (about fifteen times the mean ATT), 16 of the 22 NSNs have a negative estimate and 6 have a positive estimate, and only one of the 22 NSNs has an individually significant estimate at $p < 0.05$. The hearing-protector NSN that produces the largest outlier in the matched-controls distribution shows up here too at +1.81.

Figure F10 and Table A6 tell a different story for the aggregate event-time path. The mean effect across treated NSNs is -0.03 at EY 0, -0.07 at EY+1, -0.29 at EY+2, and -0.16 at EY+3. At EY+2 the aggregate is significantly below zero ($p < 0.01$) even though only one individual NSN’s per-NSN estimate clears the same threshold; on the natural-log scale, -0.29 corresponds to the average treated NSN’s maximum unit price being about 25% lower than its synthetic comparison at EY+2. The p-value for EY+1 exceeds the




Figure F9. Aggregate synth event-time path, domestic sourcing share. Mean effect across treated NSNs at each event year, with 95% CI band.

0.05 significance threshold, and EY+3 has too few treated NSNs (three) to test reliably.

As Figure A12 shows, mean offers produces ATTs centered close to zero. Only two of the 22 NSNs have individually significant estimates at $p < 0.05$, and these two NSNs have different signs. The aggregate event-time path (Figure F11, Table A6) is similar: the mean effect across treated NSNs oscillates between +0.15 at EY 0, -0.31 at EY+1, -0.05 at EY+2, and +0.86 at EY+3, and no event-year aggregate reaches statistical significance.

Pre-fit RMSPE measures how closely each synthetic tracks its treated NSN before the waiver – smaller is a tighter fit. Median pre-fit RMSPE is 0.0004 log points for unit price, 0.003 for domestic share, and 0.008 for mean offers (Table 3), small against the scale of each outcome.

The results come with certain caveats. First, on domestic share the median synthetic places non-trivial weight on 82 donor NSNs; the remaining donors in the eligible pool (about 14,000) receive near-zero weight (Table 3). Most candidate donors have a domestic




Figure F10. Aggregate synth event-time path, maximum logged unit price.




Figure F11. Aggregate synth event-time path, mean offers.

share of 1 across the full window, so the optimizer matches the treated NSN’s lower pre-waiver values by combining the few donors that vary. Second, the mean-offers fits face the opposite situation: smaller donor pools and noisier outcomes lead to wider placebo intervals. Third, the three-prior-event-year identification floor means this design cannot estimate effects for the five or six most recently waived NSNs.

Robustness

Three checks test whether the sourcing result depends on how the treated sample is drawn or measured. These checks are reported in the Appendix Robustness section. The first runs all three estimators on a single common sample – the most restrictive of the three, close to what the synthetic-control estimator requires – to see whether the slight differences in their samples affect the results. The picture remains the same: domestic sourcing still falls, price still falls in the long term, and offers still show no effect.

The second drops freight containers from the treated sample. Containers make up most of the treated items, so this check confirms the findings are not driven by that one product group. The remaining sample is very small: four items for synthetic controls, five or six for the other two estimators. The first-year drop in domestic sourcing is still present, but it does not persist; by the second post-waiver year the aggregate is back near zero. This reversal is in line with the intent of the waiver policy, but statistically the trend is driven by one item rather than a broad recovery. Price and offers show no reliable signal: with so few items the estimates are too noisy to read. The pooled price estimate even flips positive, but that too is driven by a single item – the hearing protector.

The third check re-estimates on data from the Defense Logistics Agency alone, dropping the supplementary BidLink transaction records that many treated items rely on. The decline in domestic sourcing holds across all three estimators, and the long-term decline in price is preserved as well; mean offers shows no effect.

In summary, the robustness checks generally agree with the main analysis. The one notable divergence is in the non-container sample, where domestic sourcing turns positive in the long term – but that result is driven by a single treated NSN.

Interpretation

Domestic sourcing share comes back negative across all three estimators at every post-waiver event year. The event study’s coefficients are negative and individually significant, but treated NSNs’ domestic share was already declining before the waiver, so its estimates absorb part of that pre-existing trend. Matched controls and synthetic controls each build comparison groups whose pre-waiver paths already mirror that decline, so the comparison accounts for the pre-existing trend, and what’s left is closer to the waiver’s own effect. Both still find that effect statistically significant.

Mean offers return post-waiver estimates that are not statistically distinguishable from zero under any of the three estimators. The event-study design also carries a pre-trend on this outcome, and FPDS leaves the offer count blank for many code-A “full and open competition” awards, so the null result is harder to read as evidence of no effect than the point estimates alone would suggest.

Maximum logged unit price moves negative under every estimator, but whether the effect is statistically significant depends on the extent of the post-period window. The event-study shows a significant reduction for event years 1 through 3. Under matched controls, restricting to event year ≥ 2 produces a significant negative effect; the full window does not. Under synthetic controls, the aggregate event-time path crosses statistical significance at event year +2 even though only one per-NSN estimate is significant individually.

The negative effect on domestic sourcing invites several readings.

The products being waived for nonavailability could be just that: truly nonavailable. Agency procurement officials, after conducting market research and securing approval from the Office of Management and Budget’s Made in America Office, have determined that the NSN is nonavailable, and thus worthy of receiving a waiver. Hence the collapse in domestic sourcing, because the product is simply not available domestically.

Potential suppliers of the waived product may not be watching the waiver portal. MadeInAmerica.gov receives roughly 0.06% as many visits as SAM.gov in a typical

30-day window (about 4,000 against 6.5 million from late April through late May 2026).¹⁴ The signal may not be reaching the people who would act on it, which is a potential explanation for the null effect on average offers received.

Even if the signal does reach potential suppliers, the supply side may not respond. This analysis is scoped to Defense procurement. The Defense Logistics Agency’s nonavailability rate is already a quarter of the federal-wide rate (0.09 percent of DLA transactions against 0.35 percent across the government), possibly driven by national-security needs. Already, defense procurement has a higher price preference for domestically-manufactured agencies than civilian agencies do¹⁵. The products easier to produce domestically may have already been reshored, leaving the nonavailable items that are harder to build in the United States. And the federal government, for all its scale, is not always a large enough buyer to justify standing up domestic production. If DLA purchases were to make up only a small share of total sales for a supplier, the cost of building US capacity could outweigh the contracting revenue at stake.

The policy is also still in its early days. The public waiver system began rolling out in late 2021. It takes time to reorient supply chains, stand up factories, and clear the federal acquisition bar – especially for Defense work. The waiver’s effect may still be working its way through the marketplace. And the treatment sample shrinks past event year 2, so the analyses are not well-positioned to detect a slow-developing response.

Waivers don’t seem to dissuade bidders (or for that matter, entice them). None of the three estimators finds a statistically significant change in average bids received. The estimators disagree on direction in the noise around zero – event study and matched controls show a short-term drop at event year 0, synth shows a small increase – but none of those movements clears significance.

Potential suppliers may not be seeing the waiver to begin with. FPDS also leaves the offer

¹⁴U.S. General Services Administration, Digital Analytics Program (U.S. General Services Administration, n.d.), 30-day window 2026-04-23 through 2026-05-23: madeinamerica.gov 4,141 visits versus sam.gov 6,462,354 visits.

¹⁵50% for the DLA, 20-30% for civilian agencies (“Subpart 225.1 - BUY AMERICAN—SUPPLIES | Acquisition.GOV,” n.d.; “Buy American Act,” n.d.)

count blank on many full-and-open competition awards, so contracts where competition might have shifted could be missing from the data entirely. And the offer count would not capture domestic suppliers who are developing capacity and preparing to enter the market, but have yet to enter a bid.

The Price hypothesis had pointed to a possible rise in unit cost: with the DFARS 50% evaluation factor in play, domestic suppliers drawn in by a posted waiver could win contracts at up to 50% above the foreign bid. But that prediction was contingent on domestic sourcing rising, which the data do not show. Since domestic sourcing decreases, the maximum unit price (logged) would be expected to decrease or remain unchanged when compared to similar non-waived NSNs. The 50% price preference offered to domestically manufactured end products would no longer apply, putting downward pressure on prices.

But if nonavailability is caused by the exit of a domestic manufacturer, then the reduction in competition would be expected to lead to an increase in prices. The unit price paid does seem to decrease in the long term. The event study sees no statistically significant change, but matched controls and synthetic controls – which account for trends preceding the waiver – see statistically significant reductions in unit price from event year 2 onwards.

What is surprising is that none of the estimators find an immediate reduction in unit prices. In event year 0, both the event study and the synthetic controls observe a change in logged maximum unit prices that is statistically indistinguishable from 0. Similarly, the matched controls analysis encompassing event year 0 onwards does not see a statistically significant reduction in logged maximum unit prices. This may indicate the gradual loss of a domestic supplier and the subsequent loss of downwards pricing pressure from competition, although there is no significant change in bids posted.

Across the three outcomes, the three estimators converge in a consistent direction: domestic sourcing falls, the offer count does not move, and the unit price falls in the long term.

Conclusion

The thesis set out to test whether publicizing nonavailability waivers changes the sourcing, competition, and pricing of the products it names. The event study, matched controls, and synthetic controls build a picture of how these waivers play out. Domestic sourcing falls. Offers per transaction do not move. The maximum logged unit price drifts down in the long term. The event study shows a pre-existing decline in domestic sourcing among the waived NSNs; matched controls and synthetic controls construct comparison paths that absorb that pre-trend, and the domestic-sourcing decline survives in both. On the supply chains this analysis can see, the waiver portal is not drawing domestic suppliers back into the products it flags – at least not yet, and not visibly within the Defense Logistics Agency.

Compared to subsidies and tariffs, the waiver portal is a relatively cheap policy instrument. Tariffs raise costs and tensions with trading partners; industrial subsidies like those of the CHIPS Act require tens of billions in appropriations. The public waiver system, in contrast, runs on a website. Its mechanism is informational: tell the public what the government cannot find domestically, so the market can fill the gap. The analysis here suggests the market does not step up, at least not on the timescale and scope available to measure. The waiver portal can be a useful audit trail and a public record of domestic supply chain gaps – the analysis here would not have been possible without it – but its capacity to move the supply side, on its own, may be more limited than was hoped when the system was established.

For information to work as a tool of industrial policy, the demand signal has to reach suppliers who would act on it. MadeInAmerica.gov receives far fewer visits than SAM.gov, and SAM is where vendors already check for federal opportunities. Out of the roughly 4,000 monthly visits MadeInAmerica.gov does receive, a sizable portion likely comes from federal officials submitting waivers to the portal, rather than companies hunting for opportunities to meet federal demand. An improvement to the waivers' reach would be to surface posted waivers on SAM directly, alongside the solicitations and contract data vendors are already filtering through. A quarterly rollup of waived items, disseminated through the

procurement channels suppliers monitor for opportunities, would lower the search cost further.

Granularity in the data itself also matters. Product Service Codes (PSC) are the most consistently available product category identifiers in waiver and procurement data, and they are too broad a unit for product-level analysis. Take PSC 6640: “Laboratory Equipment and Supplies”. This is the single most common PSC on the portal. Under that one code sit DNA extractors, femtosecond laser systems, and silicon wafers, products with very different purposes and very different manufacturing processes. A finer classification, or systematic use of NSN codes where they exist, would further illuminate supply chain gaps, allowing suppliers and government stakeholders to more easily pinpoint areas of nonavailability.

Another fix is to link the two sides of the transaction. PIIDs join waivers with their respective contracts, but the identifier is underused in both directions. Fewer than 40% of waivers in the portal record a PIID, and the contract records for procurements made under a nonavailability exemption do not in turn reference the waiver. This gap can be closed by populating the PIID field on the waiver, and disclosing the waiver ID in the corresponding contract record.

Linking waivers with contracts would make it easier to measure how much spending each waiver actually covers. A waiver record does not state the dollar value of the contract it authorizes. When PIIDs are unavailable, it becomes much more difficult to understand the degree of demand falling under a waiver, weakening the strength of the demand signal. This thesis works around the problem by measuring spend under waived NSNs instead. This approach is only possible because the DLA publishes a separate procurement dataset carrying quantity and unit price at the line-item level – detail that is rarely available otherwise, and that requires a precise product identifier most waivers do not supply. Closing the PIID gap would aid policymakers, suppliers, and researchers trying to understand just how much demand is being covered by a waiver.

Due to limitations around existing data, this analysis is scoped to the Defense Logistics

Agency – NSN coverage is densest there. But the DLA is not representative of the broader waiver population. For starters, its nonavailability rate is roughly a quarter of the federal-wide rate. Furthermore, the civilian agencies that dominate the waiver universe (like NIST and NIH) have different procurement mechanics and demand different types of products. As a result, procurement in these agencies might respond differently to public disclosure of nonavailability.

Even without a systematic change to how procurement is reported in the data, civilian agencies can be covered by assigning more granular product codes (such as NSNs) to records that do not carry them. Large language models or natural language processing could link waivers and procurement records to more precise product identifiers, using the free-text descriptions, CAGE supplier codes, and PSC codes of each record. This would make the results more generalizable, extending the analysis to the larger, civilian-dominated share of the waiver universe. It would, however, come with a potential cost to accuracy: there may not be enough information to match a record to an NSN, and the matching process can introduce errors and inconsistencies.

The portal was built on the idea that broadcasting to the public what the government cannot buy domestically will allow the market to fill the gap. On the supply chains this analysis can see, the tests do not show the desired market response. But the system is young. Rollout of the new policy began in late 2021, and the treated sample thins beyond two years after a waiver is issued. As more post-waiver years accrue, longer windows could detect slow-developing supply responses the existing data cannot. Supply chain illumination, through the public waiver portal, may yet prove to be a powerful policy lever. With changes to improve visibility, the demand signal may better reach the stakeholders it hopes to be seen by, and achieve its industrial policy goals.

Bibliography

Arkhangelsky, Dmitry, Susan Athey, David A. Hirshberg, Guido W. Imbens, and Stefan Wager. 2021. “Synthetic Difference-in-Differences.” *American Economic Review* 111 (12): 4088–118. <https://doi.org/10.1257/aer.20190159>.

Bombardini, Matilde, Andres Gonzalez-Lira, Bingjing Li, and Chiara Motta. 2024. The Increasing Cost of Buying American. No. W32953. National Bureau of Economic Research. <https://doi.org/10.3386/w32953>.

Office of Management and Budget, Made in America Office. 2025. *Building Resilience through a Made in America Industrial Strategy*. Office of Management and Budget. https://bidenwhitehouse.archives.gov/wp-content/uploads/2025/01/Building-Resilience-through-a-Made-in-America-Industrial-Strategy.final_.pdf.

Bulow, Jeremy, and Paul Klemperer. 1996. “Auctions Versus Negotiations.” *The American Economic Review* 86 (1): 180–94. JSTOR.

“Buy American Act: How Amendments Affect Canadian Businesses.” n.d. Canadian Commercial Corporation (CCC). Accessed May 24, 2026. <https://www.ccc.ca/en/insights-for-exporters/buy-american-act-how-new-amendments-affect-canadians-businesses/>.

Fed. Regist. (Federal Register). 2021. “Ensuring the Future Is Made in All of America by All of America’s Workers.” January 28. <https://www.federalregister.gov/documents/2021/01/28/2021-02038/ensuring-the-future-is-made-in-all-of-america-by-all-of-americas-workers>.

Fed. Regist. (Federal Register). 2025. “Federal Acquisition Regulation: List of Domestically Nonavailable Articles.” May 12. <https://www.federalregister.gov/documents/2025/05/12/2025-08022/federal-acquisition-regulation-list-of-domestically-nonavailable-articles>.

“Federal Awards | Advanced Search | USAspending.” n.d. Accessed May 24, 2026. <https://usaspending.gov/search>.

Ohashi, Hiroshi. 2009. “Effects of Transparency in Procurement Practices on Government Expenditure: A Case Study of Municipal Public Works.” *Review of Industrial Organization*

34 (3): 267–85. <https://doi.org/10.1007/s11151-009-9208-1>.

“Subpart 225.1 - BUY AMERICAN—SUPPLIES | Acquisition.GOV.” n.d. Accessed May 24, 2026. <https://www.acquisition.gov/dfars/subpart-225.1-buy-american%E2%80%94supplies>.

U.S. Congress. House. Committee on Veterans’ Affairs. Subcommittee on Oversight and Investigations. VA Procurement: Made in America. 118th Cong., 1st sess., September 20, 2023. Serial No. 118-32. Congress.gov. <https://www.congress.gov/event/118th-congress/house-event/LC72911/text>

U.S. General Services Administration (Analytics.Usa.Gov | The US Government’s Web Traffic.). n.d. “Analytics.Usa.Gov | The US Government’s Web Traffic.” Accessed May 23, 2026. <https://analytics.usa.gov>.

U.S. Government Accountability Office. 2018. “Buy American Act: Actions Needed to Improve Exception and Waiver Reporting and Selected Agency Guidance | U.S. GAO.” December 18. <https://www.gao.gov/products/gao-19-17>.

USAspending.gov. 2023. “Data Sources.” January 13. <https://www.usaspending.gov/data/data-sources-download.pdf>.

“Waivers | Made in America.” n.d. Accessed May 24, 2026. <https://www.madeinamerica.gov/waivers/>.

Young, Shalanda D. 2021. M-21-26 MEMORANDUM FOR HEADS OF EXECUTIVE DEPARTMENTS AND AGENCIES. June 11.

Appendix

A Primer on BAA and BABA

The online Made in America waiver portal contains different types of waivers, some of which are outside the scope of this thesis. Broadly, waivers fall under two laws: Buy American Act (BAA), and Build America Buy America (BABA). Buy American is the longstanding policy that imposes domestic manufacturing requirements on government procurement. BABA is a relatively recent regulation that applies similar requirements to infrastructure projects funded by the federal government.

Because BABA requirements apply to grants to non-federal organizations, there is no standardized dataset of procurement by these organizations, like USAspending provides for the federal government. As a result, it would not be possible to compare how the origin of manufactured goods changes over time for these third-parties.

That leaves Buy American waivers. The Made in America website provides two datasets for procurement waivers: nonavailability, and urgent requirement reports. Because the primary goal of this thesis is to examine how public waivers have affected domestic non-availability, nonavailability waivers are the focus. Urgent requirement waivers cover narrow, time-sensitive procurements rather than ongoing sourcing gaps and fall outside this thesis's scope.

Exceptions to the Exception

Not all contracts that are exempted for nonavailability need public waivers. Under both the current and prior systems, contracts that meet all 3 of the following requirements do not need a “written determination”: sourced under full and open competition, properly synopsisized (advertised to the public), and no offer for a domestic end product was received (FAR 25.103(b)(3)). In other words, if private industry had a fair chance to offer a domestically manufactured product and no such offer came, a written determination (i.e. a public waiver) is not needed.

This fits with the illumination goal of the public waivers. If private industry might not have had adequate opportunity to respond to the government need and subsequent domestic shortcoming, public waivers could put a spotlight on the nonavailability gap.

Additionally, FAR maintains a list of nonavailable articles for which a public waiver is not needed (unless a satisfactory domestic bid appears during the contracting phase)¹⁶. This list mainly consists of raw materials like corned beef, olives, and yeast. There are a few articles that stand out: spare parts for foreign equipment (that don’t have domestically-manufactured parts) and books or other literary material that aren’t otherwise printed in the US.

¹⁶The list has since been substantially shortened: a rule proposed in October 2024 and finalized June 2025 removed nearly two-thirds of its 109 articles. OMB has projected this would be the shortest the nonavailability list has ever been. (Fed. Regist. 2025; Office of Management and Budget, Made in America Office 2025)

Foreign Manufacture Share




Figure A0. Nonavailability share of foreign-manufacture transactions by fiscal year.

Treated NSN Summary

Table A1. Treated NSNs: Product, Waiver Date, and Procurement Summary

NSN	Item name	First waiver	Transactions	Avg unit price	Avg bids
8150-01-463-8553	Container, freight, utility	2022-07-06	267	\$6,002	4.9
8150-01-574-1682	Container, freight, utility	2024-03-28	213	\$6,096	2.7
8150-01-533-8673	Tricon, freight, specific purpose	2024-03-28	195	\$6,542	2.9
8150-01-501-2741	Container, freight, utility	2022-12-21	160	\$8,287	2.5
8150-01-463-3177	Container, freight, utility	2022-09-28	137	\$3,979	3.4
8150-01-483-9123	Container, freight, utility	2022-10-17	136	\$6,005	2.8
8150-01-533-8674	Tricon, freight, specific purpose	2022-11-08	135	\$6,130	3.3
8150-01-533-8677	Quadcon, freight, specific purpose	2024-03-28	127	\$6,096	3.1
8150-01-481-3741	Bicon, freight, general purpose	2022-10-17	96	\$6,750	2.6
8150-01-463-8555	Container, freight, utility	2022-08-25	87	\$8,758	3.2
8150-01-483-9125	Container, freight, utility	2024-03-28	86	\$6,423	2.5
8150-01-491-2678	Container, freight, utility	2022-09-27	81	\$7,554	4.1
8150-01-533-8675	Quadcon, freight, specific purpose	2023-03-21	66	\$5,478	2.7
8150-01-514-8414	Container, freight, utility	2022-10-17	64	\$9,137	2.8
8150-01-483-9115	Container, freight, utility	2023-03-22	58	\$5,565	2.7
8150-01-483-9121	Container, freight, utility	2022-10-17	57	\$5,410	2.5
8150-01-513-4913	Container, freight, utility	2024-03-28	49	\$11,086	2.9
8150-01-481-3740	Bicon, freight, general purpose	2024-03-28	44	\$7,348	2.4
8150-01-488-6545	Container, freight, utility	2023-01-19	35	\$9,589	2.8
3990-01-574-2050	Roller, material handling	2023-04-13	32	\$1,599	2.4
8150-01-573-7424	Container, freight, utility	2022-09-22	24	\$7,634	2.6
4510-01-527-2274	Shower bath fixture	2023-03-08	21	\$362	14.0
4240-01-538-7970	Protector, hearing	2023-03-06	16	\$37	12.1
8150-01-491-2727	Container, freight, utility	2024-02-22	10	\$13,751	2.6
1680-01-622-9189	Enhanced ground proximity warning system	2022-09-13	6	—	1.0
1660-01-600-4909	Heat exchanger, air-to-air, aircraft	2022-11-22	4	\$13,722	1.0
2895-14-554-3146	Motor, pneumatic	2023-03-14	4	\$4,070	13.2
Average across NSNs	21 containers, 6 items	—	81.9	\$6,670	3.9

The 27 treated NSNs used in the analysis (21 freight containers and six individual items), sorted by transaction count. Transactions are counted over fiscal years 2017 through January 2026 and total 2,210. Avg unit price is the mean unit price observed for the NSN; for 1660-01-600-4909 (heat exchanger) it is a proxy recovered as contract obligation $\div$ quantity, and for 1680-01-622-9189 (excluded from the unit-price analysis) it is shown as a dash. Avg bids is the mean number of offers received per transaction. The final row reports the average across the 27 NSNs (unit price excludes the one NSN without a recorded price).

Event Study Specification Ladder

The tables and figures below report the event-study coefficients under successively richer specifications: pooled (no fixed effects), NSN fixed effects only, NSN and FY fixed effects, and the full specification with NSN + FY fixed effects and standard errors clustered on NSN. The full-spec column matches the event-study results presented in the body.

Table A2. Event-study coefficient ladder, domestic sourcing share.

Event year	(1) OLS	(2) +NSN FE	(3) +NSN+FY FE
$k = -8$	+0.677*** (0.073)	+0.857*** (0.056)	+0.826*** (0.056)
$k = -7$	+0.658*** (0.074)	+0.763*** (0.077)	+0.735*** (0.077)
$k = -6$	+0.643*** (0.076)	+0.649*** (0.071)	+0.625*** (0.071)
$k = -5$	+0.539*** (0.068)	+0.547*** (0.054)	+0.530*** (0.054)
$k = -4$	+0.353*** (0.115)	+0.374*** (0.106)	+0.362*** (0.105)
$k = -3$	+0.161* (0.086)	+0.168** (0.081)	+0.164** (0.081)
$k = -2$	+0.182* (0.094)	+0.162* (0.096)	+0.161* (0.095)
$k = -1$	<i>reference</i>		
$k = +0$	-0.297*** (0.075)	-0.286*** (0.072)	-0.285*** (0.073)
$k = +1$	-0.222*** (0.082)	-0.216*** (0.068)	-0.213*** (0.068)
$k = +2$	-0.200** (0.094)	-0.243*** (0.089)	-0.237*** (0.089)
$k = +3$	-0.229*** (0.082)	-0.324*** (0.065)	-0.314*** (0.065)
NSN FE	—	✓	✓
FY FE	—	—	✓
Joint pre-trend p	0.000***	0.000***	0.000***
Avg post β ($k \geq 0$)	-0.237*** (0.074)	-0.267*** (0.064)	-0.262*** (0.065)
N obs	1,468,681	1,468,681	1,468,681
R ²	0.001	0.674	0.676

Coefficients on event-year dummies. Standard errors in parentheses, clustered on NSN. Reference event year is $k = -1$. Joint pre-trend p is the Wald F-test that all pre-treatment dummies are zero. Avg post β is the linear combination of post-period coefficients ($k \geq 0$) with its CRV1 standard error. Specification (3), with NSN and FY fixed effects, is the one reported in the main Analysis section. Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A3. Event-study coefficient ladder, maximum logged unit price.

Event year	(1) OLS	(2) +NSN FE	(3) +NSN+FY FE
$k = -8$	+0.023 (0.225)	-0.196 (0.186)	+0.118 (0.183)
$k = -7$	-0.980* (0.545)	-0.301*** (0.113)	-0.008 (0.115)
$k = -6$	-0.885** (0.442)	-0.260*** (0.073)	-0.016 (0.081)
$k = -5$	-0.510** (0.240)	-0.195*** (0.068)	+0.023 (0.075)
$k = -4$	-0.435* (0.223)	-0.202*** (0.076)	-0.025 (0.083)
$k = -3$	-0.643 (0.404)	-0.176* (0.093)	-0.044 (0.100)
$k = -2$	-0.341 (0.249)	-0.006 (0.055)	+0.073 (0.059)
$k = -1$	<i>reference</i>		
$k = +0$	-0.241 (0.159)	+0.008 (0.080)	-0.077 (0.082)
$k = +1$	-0.123 (0.186)	+0.013 (0.047)	-0.150*** (0.053)
$k = +2$	-0.196 (0.183)	+0.047 (0.060)	-0.187*** (0.063)
$k = +3$	+0.022 (0.157)	-0.041 (0.106)	-0.302*** (0.116)
NSN FE	—	✓	✓
FY FE	—	—	✓
Joint pre-trend p	0.187	0.004***	0.462
Avg post β ($k \geq 0$)	-0.134 (0.123)	+0.007 (0.057)	-0.179*** (0.062)
N obs	1,468,457	1,468,457	1,468,457
R ²	0.000	0.975	0.978

Coefficients on event-year dummies. Standard errors in parentheses, clustered on NSN. Reference event year is $k = -1$. Joint pre-trend p is the Wald F-test that all pre-treatment dummies are zero. Avg post β is the linear combination of post-period coefficients ($k \geq 0$) with its CRV1 standard error. Specification (3), with NSN and FY fixed effects, is the one reported in the main Analysis section.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A4. Event-study coefficient ladder, mean offers.

Event year	(1) OLS	(2) +NSN FE	(3) +NSN+FY FE
$k = -8$	-0.082 (0.446)	+0.214 (0.317)	-1.897*** (0.316)
$k = -7$	+0.685 (1.166)	-0.489* (0.273)	-2.424*** (0.276)
$k = -6$	+1.775 (1.251)	+0.671 (0.504)	-0.867* (0.491)
$k = -5$	+1.700 (1.576)	+0.708 (1.237)	-0.262 (1.239)
$k = -4$	+0.190 (0.614)	-0.660 (0.590)	-1.032* (0.596)
$k = -3$	+1.012 (0.736)	+0.091 (0.200)	-0.212 (0.201)
$k = -2$	+0.275 (0.616)	-0.473* (0.284)	-0.659** (0.281)
$k = -1$	<i>reference</i>		
$k = +0$	+0.128 (0.532)	-0.875 (0.711)	-0.791 (0.703)
$k = +1$	-0.081 (0.524)	-0.760* (0.407)	-0.688* (0.394)
$k = +2$	+0.286 (0.656)	-0.602* (0.345)	-0.534* (0.323)
$k = +3$	+0.546 (0.493)	-0.307 (0.863)	-0.237 (0.805)
NSN FE	—	✓	✓
FY FE	—	—	✓
Joint pre-trend p	0.067*	0.000***	0.000***
Avg post β ($k \geq 0$)	+0.220 (0.458)	-0.636 (0.454)	-0.562 (0.432)
N obs	1,120,312	1,120,312	1,120,312
R ²	0.000	0.553	0.570

Coefficients on event-year dummies. Standard errors in parentheses, clustered on NSN. Reference event year is $k = -1$. Joint pre-trend p is the Wald F-test that all pre-treatment dummies are zero. Avg post β is the linear combination of post-period coefficients ($k \geq 0$) with its CRV1 standard error. Specification (3), with NSN and FY fixed effects, is the one reported in the main Analysis section.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.




Figure A1. Event-study coefficient ladder (fixed effects added one at a time), domestic sourcing share.




Figure A2. Event-study coefficient ladder (fixed effects added one at a time), maximum logged unit price.




Figure A3. Event-study coefficient ladder (fixed effects added one at a time), mean offers.

Matched Controls: Per-NSN Distributions

The dotplots below show each per-NSN τ_i for the matched-controls design, by outcome and post-period window. The vertical band marks the 95% confidence interval for the ATT. The main ATTs and same-FSG sensitivity values appear in Table 2 of the body.




Figure A4. Per-NSN τ_i distribution with ATT 95% CI band, domestic sourcing share ($EY \geq 0$).




Figure A5. Per-NSN τ_i distribution with ATT 95% CI band, domestic sourcing share ($EY \geq 2$ window).




Figure A6. Per-NSN τ_i distribution with ATT 95% CI band, maximum logged unit price (EY ≥ 0).




Figure A7. Per-NSN τ_i distribution with ATT 95% CI band, maximum logged unit price ($EY \geq 2$ window).




Figure A8. Per-NSN τ_i distribution with ATT 95% CI band, mean offers ($EY \geq 0$).




Figure A9. Per-NSN τ_i distribution with ATT 95% CI band, mean offers ($EY \geq 2$ window).

Synthetic Controls: Per-NSN Detail

The forest plots below show each per-NSN synthdid ATT with its placebo 95% CI. Tables A5 and A6 report the per-NSN ATT distribution and aggregate event-time path, respectively.




Figure A10. Per-NSN synth ATT with placebo 95% CIs, domestic sourcing share.

Table A5. Synthetic Controls: Per-NSN ATT Distribution by Outcome.

Outcome	N fits	Median ATT	Mean ATT	SD across NSNs	Significant ($p < 0.05$)
Domestic sourcing share	23	-0.414	-0.443	0.267	21 / 23
Maximum logged unit price	22	-0.116	-0.030	0.443	1 / 22
Mean offers	22	+0.033	+0.282	2.176	2 / 22

ATTs are estimated per treated NSN using `synthdid` (Arkhangelsky et al. 2021). Significance is assessed against the random-subsample placebo distribution (200 placebos per fit). Pre-fit RMSPE is the root mean squared prediction error in the pre-waiver period — smaller means the synthetic tracks its treated NSN more closely before the waiver. The “active donors / pool size” column reports the median count of non-treated NSNs that receive non-trivial weight ($w > 1 / (2 \cdot \text{pool size})$) against the total eligible pool — lower ratios indicate that the synthetic relies on a small subset of the pool.




Figure A11. Per-NSN synth ATT, maximum logged unit price.




Figure A12. Per-NSN synth ATT, mean offers.

Table A6. Synthetic Controls: Aggregate Event-Time Path by Outcome.

Outcome	Event year	N NSNs	Mean effect	SE
Domestic sourcing share	+0	23	-0.461***	0.059
	+1	20	-0.367***	0.052
	+2	15	-0.415***	0.092
	+3	3	-0.682**	0.105
Maximum logged unit price	+0	22	-0.025	0.100
	+1	20	-0.068	0.041
	+2	14	-0.291***	0.050
	+3	3	-0.162	0.180
Mean offers	+0	22	+0.153	0.515
	+1	20	-0.311	0.314
	+2	14	-0.051	0.238
	+3	3	+0.863	0.322

Aggregate effect at each post-waiver event year is the mean of per-NSN demeaned effects (treated minus synthetic, with the pre-waiver mean subtracted to remove any constant fit gap). SE and significance stars are from a one-sample t-test on the per-NSN effects at that event year ($df = N - 1$). Pre-waiver event years are flat by construction and not shown.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Robustness

This section reports three checks on the treatment-effect results. Each re-estimates the event study, matched controls, and synthetic controls, then sets the result beside the main estimates. The first holds the treated sample fixed across the three estimators. The second removes the product group that dominates the treated items. The third re-estimates on data from the Defense Logistics Agency alone, leaving out the supplementary BidLink transactions that extend coverage for many treated items but reach only a few of the controls.

Common Treated Sample

In the Analysis section, the three estimators do not run on identical treated samples. Each imposes its own data requirements, and the synthetic-control estimator is the most demanding, dropping any item observed in fewer than three pre-waiver event years (so that there is enough pre-waiver history to build a credible counterfactual). The event study and matched controls do not impose that requirement, so they retain some items the synthetic-control estimator drops. The main estimates are therefore not, strictly, comparing the same items. This check runs all three estimators on the common most-restrictive sample of 22 items.¹⁷

On the common sample, the estimates stay close to their main-analysis values. For domestic sourcing, the event study holds at -0.26 and synthetic controls at -0.44 (against -0.46), while matched controls give -0.38 over the full post-period and -0.23 at event year ≥ 2 (against -0.40 and -0.23). The long-term decline in price carries over as well: the matched-controls effect at event year ≥ 2 is -0.17 ($p < 0.05$) and the synthetic-control event-time path again turns significantly negative at event year 2 (-0.29, $p < 0.01$), both close to the main analysis values, while the event study's post-period price path stays negative (average -0.17). Offers show no effect in any estimator, same as in the main analysis. The results from the estimators thus survive using the uniform, more restrictive treatment sample.

¹⁷One waived item, an enhanced ground proximity warning system, has no recorded unit prices. Matched controls draw their matches from each item's pre-waiver price, sourcing, and volume, so they cannot match this product and exclude it from every outcome. To keep one sample common to all three estimators, it is dropped throughout this check.

Table A7. Common Sample: Robustness Estimates vs the Main Analysis

Estimator	Outcome	Main analysis	Common Sample
Event study (avg. post-period coefficient)	Domestic sourcing share	-0.262	-0.263
	Maximum logged unit price	-0.179	-0.171
	Mean offers	-0.562	-0.435
Matched controls, ATT (event year ≥ 0)	Domestic sourcing share	-0.397	-0.380
	Maximum logged unit price	-0.117	-0.053
	Mean offers	-0.494	-0.096
Matched controls, ATT (event year ≥ 2)	Domestic sourcing share	-0.226	-0.226
	Maximum logged unit price	-0.170	-0.170
	Mean offers	-0.822	-0.822
Synthetic controls (mean per-NSN ATT)	Domestic sourcing share	-0.443	-0.441
	Maximum logged unit price	-0.030	-0.030
	Mean offers	+0.282	+0.282

Each robustness estimate beside its main-analysis (full-sample) value for the same estimator. The event study reports the average post-period event-year coefficient; matched controls the collapsed pre/post ATT in two post-period windows (all post-waiver years, and event year ≥ 2 onward); synthetic controls the mean per-NSN ATT.

Table A8. Common Sample: Event-Study Coefficient by Event Year and Outcome (vs Main Analysis)

Event year	Domestic sourcing share		Maximum logged unit price		Mean offers	
	Common Sample	Main	Common Sample	Main	Common Sample	Main
-4	+0.388*** (0.107)	+0.362***	-0.063 (0.084)	-0.025	-0.511 (0.470)	-1.032*
-3	+0.164** (0.082)	+0.164**	-0.044 (0.102)	-0.044	-0.189 (0.205)	-0.212
-2	+0.168* (0.099)	+0.161*	+0.073 (0.059)	+0.073	-0.648** (0.281)	-0.659**
-1	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)
0	-0.259*** (0.072)	-0.285***	-0.056 (0.085)	-0.077	-0.228 (0.374)	-0.791
+1	-0.227*** (0.069)	-0.213***	-0.139*** (0.053)	-0.150***	-0.729* (0.406)	-0.688*
+2	-0.257*** (0.091)	-0.237***	-0.187*** (0.063)	-0.187***	-0.540* (0.318)	-0.534*
+3	-0.310*** (0.067)	-0.314***	-0.302*** (0.116)	-0.302***	-0.244 (0.804)	-0.237

Event-year coefficient for the Common Sample beside the main-analysis value, by outcome. Common Sample column: coefficient with CRV1 SE clustered on NSN in parentheses; main column: coefficient only. NSN + FY fixed effects; the reference period is event year -1.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A9. Common Sample: Matched-Controls ATT by Outcome and Window

Outcome	Post window	N	Main matching ATT (SE)	Same-FSG sensitivity ATT (SE)
Domestic sourcing share	event year ≥ 0	22	-0.380*** (0.076)	-0.310*** (0.076)
	event year ≥ 2	14	-0.226* (0.111)	-0.271** (0.101)
Maximum logged unit price	event year ≥ 0	22	-0.053 (0.095)	-0.087 (0.109)
	event year ≥ 2	14	-0.170** (0.058)	-0.224*** (0.045)
Mean offers	event year ≥ 0	22	-0.096 (0.600)	+0.432 (0.420)
	event year ≥ 2	13	-0.822 (0.486)	-0.007 (0.317)

ATT is the mean of the per-NSN τ_i , with HC1 (heteroskedasticity-robust) standard errors. *Main matching* picks the three nearest donors by Mahalanobis distance from the full eligible pool. *Same-FSG sensitivity* adds an exact-match constraint requiring each donor to share the treated NSN's two-digit Federal Supply Group (the first two digits of the four-digit FSC code).

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A10. Common Sample: Matched-Controls Fit Diagnostics (Covariate Balance)

Match-quality statistic	Main analysis	Common Sample
Mean SMD , before matching	1.386	1.358
Mean SMD , after matching	0.207	0.222
Max SMD , after matching	0.950	0.950
Share balanced ($ SMD < 0.1$), after matching	32%	27%

Standardized mean differences between each treated NSN and its matched control pool, averaged over treated NSNs and the four matching covariates (the propensity distance is excluded). Lower is better balance; $|SMD| < 0.1$ is the conventional threshold. Matching is invariant to the treated set, so these are the main-analysis matches restricted to this check's items.

Table A11. Common Sample: Synthetic Controls – Event-Time Path by Year (vs Main Analysis)

Outcome	Event year	Common Sample N	Common Sample mean	SE	Main N	Main mean
Domestic sourcing share	+0	22	-0.436***	0.056	23	-0.461***
	+1	20	-0.367***	0.052	20	-0.367***
	+2	14	-0.445***	0.093	15	-0.415***
	+3	3	-0.682**	0.105	3	-0.682**
Maximum logged unit price	+0	22	-0.025	0.100	22	-0.025
	+1	20	-0.068	0.041	20	-0.068
	+2	14	-0.291***	0.050	14	-0.291***
	+3	3	-0.162	0.180	3	-0.162
Mean offers	+0	22	+0.153	0.515	22	+0.153
	+1	20	-0.311	0.314	20	-0.311
	+2	14	-0.051	0.238	14	-0.051
	+3	3	+0.863	0.322	3	+0.863

Aggregate synthetic-control effect (mean of per-NSN demeaned effects) at each post-waiver event year, for the Common Sample beside the main-analysis aggregate, each with its own fit count N. SE and significance stars from a one-sample t-test on the per-NSN effects ($df = N - 1$); both mean columns carry stars. Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A12. Common Sample: Synthetic Controls – Per-NSN ATT Distribution

Outcome	N fits	Median ATT	Mean ATT	SD across NSNs	Significant ($p < 0.05$)	Median pre-fit RMSPE	Median active donors / pool size
Domestic sourcing share	22	-0.398	-0.441	0.273	20 / 22	0.0033	71 / 13,965
Maximum logged unit price	22	-0.116	-0.030	0.443	1 / 22	0.0004	4,730 / 15,568
Mean offers	22	+0.033	+0.282	2.176	2 / 22	0.0076	1,877 / 3,057

ATTs are estimated per treated NSN using `synthdid` (Arkhangelsky et al. 2021). Significance is assessed against the random-subsample placebo distribution (200 placebos per fit). Pre-fit RMSPE is the root mean squared prediction error in the pre-waiver period – smaller means the synthetic tracks its treated NSN more closely before the waiver. The “active donors / pool size” column reports the median count of non-treated NSNs that receive non-trivial weight ($w > 1 / (2 \cdot \text{pool size})$) against the total eligible pool – lower ratios indicate that the synthetic relies on a small subset of the pool.

Excluding Freight Containers

Freight containers account for 21 of the 27 treated items, so there is a risk that the pooled estimates are being driven by a single product group. This check drops the containers and re-estimates on the items that remain. The remaining sample is small – at most six items, and as few as three for the price and offers outcomes under the synthetic-control estimator – so these estimates are imprecise, and one item can move a pooled figure.

The immediate decline in domestic sourcing survives without the containers. In the waiver year the event-study coefficient is -0.52 and the synthetic-control effect -0.67, both sharply negative; matched controls, which pool the post-period into a single effect, give -0.37, close to the full-sample value of -0.40.

In the long-term window, the trend in domestic sourcing turns positive. By the second post-waiver year the event-study and synthetic-control paths are near or above zero (event study +0.20, synthetic controls +0.12), though neither second-year estimate is statistically significant, and the matched-controls effect at event year ≥ 2 is +0.36. Taken at face value, this shows domestic sourcing moving in the way the public waiver system intends – a disclosed supply gap drawing production back home.

But the results are heavily flawed. The long-term aggregate is driven by one item: the shower bath fixture, whose domestic-sourcing effect at event year ≥ 2 is +0.58 in synthetic controls and +1.02 in matched controls. The other items observed that far out do not follow it: the materials-handling roller is still declining, and the ground-proximity warning system, sharply negative in the waiver year, recovers to roughly zero. The remaining three items drop out of the sample before the long-term window (Table A19). With one item pulling the average up and the next pushing it down, on a two-to-four item sample, the long-term increase does not indicate a market response. The Figure A13 shows the lone item set apart from an otherwise significantly negative group.

Price and offers suffer from the same shallow sample. Mean offers shows no effect, as in the main analysis. The unit price estimate turns positive in all three estimators, but that is due to a single hearing-protection item whose price rose sharply after its waiver, set apart

from the others in Figure A14.




Figure A13. Synthetic controls, per-NSN treatment effect on domestic sourcing share (non-container treated items).

Table A13. Non-Container: Robustness Estimates vs the Main Analysis

Estimator	Outcome	Main analysis	Non-Container
Event study (avg. post-period coefficient)	Domestic sourcing share	-0.262	-0.092
	Maximum logged unit price	-0.179	+0.331
	Mean offers	-0.562	-1.594
Matched controls, ATT (event year ≥ 0)	Domestic sourcing share	-0.397	-0.373
	Maximum logged unit price	-0.117	+0.117
	Mean offers	-0.494	-1.245
Matched controls, ATT (event year ≥ 2)	Domestic sourcing share	-0.226	+0.355
	Maximum logged unit price	-0.170	-0.288
	Mean offers	-0.822	-2.689
Synthetic controls (mean per-NSN ATT)	Domestic sourcing share	-0.443	-0.415
	Maximum logged unit price	-0.030	+0.551
	Mean offers	+0.282	+1.273

Each robustness estimate beside its main-analysis (full-sample) value for the same estimator. The event study reports the average post-period event-year coefficient; matched controls the collapsed pre/post ATT in two post-period windows (all post-waiver years, and event year ≥ 2 onward); synthetic controls the mean per-NSN ATT.




Figure A14. Synthetic controls, per-NSN treatment effect on maximum logged unit price (non-container treated items).

Table A14. Non-Container: Event-Study Coefficient by Event Year and Outcome (vs Main Analysis)

Event year	Domestic sourcing share		Maximum logged unit price		Mean offers	
	Non-Container	Main	Non-Container	Main	Non-Container	Main
-4	+0.212 (0.281)	+0.362***	+0.657*** (0.082)	-0.025	-1.778 (2.097)	-1.032*
-3	+0.121 (0.133)	+0.164**	+0.065 (0.251)	-0.044	+0.768 (0.660)	-0.212
-2	+0.231 (0.186)	+0.161*	+0.322*** (0.044)	+0.073	+1.654*** (0.363)	-0.659**
-1	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)
0	-0.515* (0.284)	-0.285***	+0.571 (0.375)	-0.077	-2.549 (3.123)	-0.791
+1	+0.041 (0.137)	-0.213***	+0.281*** (0.067)	-0.150***	-1.524 (1.323)	-0.688*
+2	+0.198 (0.167)	-0.237***	+0.141 (0.114)	-0.187***	-0.709*** (0.199)	-0.534*

Event-year coefficient for the Non-Container sample beside the main-analysis value, by outcome.
 Non-Container column: coefficient with CRV1 SE clustered on NSN in parentheses; main column: coefficient only. NSN + FY fixed effects; the reference period is event year -1.
 Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A15. Non-Container: Matched-Controls ATT by Outcome and Window

Outcome	Post window	N	Main matching ATT (SE)	Same-FSG sensitivity ATT (SE)
Domestic sourcing share	event year ≥ 0	5	-0.373 (0.278)	-0.422 (0.219)
	event year ≥ 2	2	+0.355 (0.662)	+0.111 (0.489)
Maximum logged unit price	event year ≥ 0	5	+0.117 (0.420)	+0.176 (0.477)
	event year ≥ 2	2	-0.288 (0.246)	-0.349 (0.072)
Mean offers	event year ≥ 0	5	-1.245 (3.676)	-1.440 (3.728)
	event year ≥ 2	1	-2.689 (—)	+0.690 (2.403)

ATT is the mean of the per-NSN τ_i , with HC1 (heteroskedasticity-robust) standard errors. *Main matching* picks the three nearest donors by Mahalanobis distance from the full eligible pool. *Same-FSG sensitivity* adds an exact-match constraint requiring each donor to share the treated NSN’s two-digit Federal Supply Group (the first two digits of the four-digit FSC code).

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A16. Non-Container: Matched-Controls Fit Diagnostics (Covariate Balance)

Match-quality statistic	Main analysis	Non-Container
Mean SMD , before matching	1.386	1.468
Mean SMD , after matching	0.207	0.107
Max SMD , after matching	0.950	0.409
Share balanced ($ SMD < 0.1$), after matching	32%	60%

Standardized mean differences between each treated NSN and its matched control pool, averaged over treated NSNs and the four matching covariates (the propensity distance is excluded). Lower is better balance; $|SMD| < 0.1$ is the conventional threshold. Matching is invariant to the treated set, so these are the main-analysis matches restricted to this check’s items.

Table A17. Non-Container: Synthetic Controls – Event-Time Path by Year (vs Main Analysis)

Outcome	Event year	Non-Container N	Non-Container mean	SE	Main N	Main mean
Domestic sourcing share	+0	4	-0.668**	0.198	23	-0.461***
	+1	2	-0.083	0.167	20	-0.367***
	+2	3	+0.121	0.239	15	-0.415***
Maximum logged unit price	+0	3	+0.579	0.616	22	-0.025
	+1	2	-0.069	0.272	20	-0.068
	+2	2	-0.127	0.055	14	-0.291***
Mean offers	+0	3	+0.708	4.317	22	+0.153
	+1	2	-2.415	2.923	20	-0.311
	+2	2	-1.501	0.804	14	-0.051

Aggregate synthetic-control effect (mean of per-NSN demeaned effects) at each post-waiver event year, for the Non-Container sample beside the main-analysis aggregate, each with its own fit count N. SE and significance stars from a one-sample t-test on the per-NSN effects ($df = N - 1$); both mean columns carry stars.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A18. Non-Container: Synthetic Controls – Per-NSN ATT Distribution

Outcome	N fits	Median ATT	Mean ATT	SD across NSNs	Significant ($p < 0.05$)	Median pre-fit RMSPE	Median active donors / pool size
Domestic sourcing share	4	-0.366	-0.415	0.455	3 / 4	0.0002	5,056 / 18,722
Maximum logged unit price	3	+0.068	+0.551	1.096	1 / 3	0.0000	8,089 / 16,274
Mean offers	3	-0.135	+1.273	6.724	2 / 3	0.0125	1,297 / 3,237

ATTs are estimated per treated NSN using `synthdid` (Arkhangelsky et al. 2021). Significance is assessed against the random-subsample placebo distribution (200 placebos per fit). Pre-fit RMSPE is the root mean squared prediction error in the pre-waiver period – smaller means the synthetic tracks its treated NSN more closely before the waiver. The “active donors / pool size” column reports the median count of non-treated NSNs that receive non-trivial weight ($w > 1 / (2 \cdot \text{pool size})$) against the total eligible pool – lower ratios indicate that the synthetic relies on a small subset of the pool.

Table A19. Non-Container: Per-Item Domestic-Sourcing Effect

Item	Synthetic controls (effect by event year)		Matched controls (τ by event-year window)		
	0	1	2	≥ 0	≥ 2
Heat exchanger	–	–	–	+0.000	–
Ground-proximity warning system	-1.003	–	+0.002	–	–
Pneumatic motor	–	–	–	-1.000	–
Materials-handling roller	-0.223	-0.250	-0.221	-0.273	-0.307
Hearing protector	-1.000	–	–	-1.000	–
Shower bath fixture	-0.447	+0.084	+0.581	+0.406	+1.018

Domestic-sourcing effect per non-container item. Synthetic-control columns are the demeaned effect (treated minus synthetic) at each post-waiver event year; matched-controls columns are the per-item τ over all post-waiver years (event year ≥ 0) and from event year ≥ 2 onward. A dash means the estimator produced no estimate for that item or window (item dropped, or no transactions in that window).

Using Defense Logistics Agency Data Only

The transaction record for each item combines two sources: contract data from the Defense Logistics Agency, which covers every item in the analysis, and transactions recorded by BidLink, which extend that coverage for a subset of items. Many treated items draw on BidLink-recorded transactions, and for several of them, those transactions are what give the item enough activity before and after its waiver to measure at all. This check runs all three estimators on the DLA data alone, so the treated and control items are drawn from the same source.

Using the DLA data alone leaves 23 treated items in the event study and matched controls, and 21 in synthetic controls. The items that depend on the supplementary transactions fall out, having too little DLA activity around their waivers to measure. The domestic-sourcing decline holds in all three, in both the immediate and the long-term window. The event-study path remains significantly negative across the post-period, averaging -0.29 (versus -0.26 in the main analysis). Matched controls give -0.46 over the full post-period and -0.31 at event year ≥ 2 (against -0.40 and -0.23), both still significant ($p < 0.05$). The synthetic-control event-time path remains significantly negative at every post-waiver year (-0.44 at the waiver year, -0.45 at event year 2), significant for 19 of the 21 fitted items.

The long-term decline in price also holds: matched controls give -0.20 at event year ≥ 2 and -0.13 over the full post-period, and the synthetic-control event-time path again turns significantly negative at event year 2 (-0.29, $p < 0.01$), while the event study's price path stays negative (-0.18). Mean offers again shows no effect except at event year +3, which relies on only three treated NSNs.

Table A20. DLA-Only: Robustness Estimates vs the Main Analysis

Estimator	Outcome	Main analysis	DLA-Only
Event study (avg. post-period coefficient)	Domestic sourcing share	-0.262	-0.289
	Maximum logged unit price	-0.179	-0.183
	Mean offers	-0.562	-0.122
Matched controls, ATT (event year ≥ 0)	Domestic sourcing share	-0.397	-0.460
	Maximum logged unit price	-0.117	-0.133
	Mean offers	-0.494	-0.367
Matched controls, ATT (event year ≥ 2)	Domestic sourcing share	-0.226	-0.308
	Maximum logged unit price	-0.170	-0.198
	Mean offers	-0.822	-0.658
Synthetic controls (mean per-NSN ATT)	Domestic sourcing share	-0.443	-0.452
	Maximum logged unit price	-0.030	-0.078
	Mean offers	+0.282	+0.320

Each robustness estimate beside its main-analysis (full-sample) value for the same estimator. The event study reports the average post-period event-year coefficient; matched controls the collapsed pre/post ATT in two post-period windows (all post-waiver years, and event year ≥ 2 onward); synthetic controls the mean per-NSN ATT.

Table A21. DLA-Only: Event-Study Coefficient by Event Year and Outcome (vs Main Analysis)

Event year	Domestic sourcing share		Maximum logged unit price		Mean offers	
	DLA-Only	Main	DLA-Only	Main	DLA-Only	Main
-4	+0.411*** (0.118)	+0.362***	-0.018 (0.088)	-0.025	-0.897 (0.660)	-1.032*
-3	+0.131 (0.094)	+0.164**	-0.052 (0.103)	-0.044	+0.016 (0.264)	-0.212
-2	+0.116 (0.102)	+0.161*	+0.067 (0.057)	+0.073	-0.349 (0.267)	-0.659**
-1	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)	0 (ref.)
0	-0.273*** (0.081)	-0.285***	-0.117* (0.064)	-0.077	-0.407 (0.564)	-0.791
+1	-0.269*** (0.074)	-0.213***	-0.137** (0.055)	-0.150***	-0.327 (0.298)	-0.688*
+2	-0.279*** (0.101)	-0.237***	-0.186*** (0.066)	-0.187***	-0.133 (0.263)	-0.534*
+3	-0.334*** (0.077)	-0.314***	-0.291*** (0.111)	-0.302***	+0.379 (0.363)	-0.237

Event-year coefficient for the DLA-Only sample beside the main-analysis value, by outcome.

DLA-Only column: coefficient with CRV1 SE clustered on NSN in parentheses; main column: coefficient only. NSN + FY fixed effects; the reference period is event year -1.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A22. DLA-Only: Matched-Controls ATT by Outcome and Window

Outcome	Post window	N	Main matching ATT (SE)	Same-FSG sensitivity ATT (SE)
Domestic sourcing share	event year ≥ 0	23	-0.460*** (0.074)	-0.371*** (0.076)
	event year ≥ 2	13	-0.308** (0.129)	-0.328** (0.111)
Maximum logged unit price	event year ≥ 0	23	-0.133** (0.052)	-0.190*** (0.048)
	event year ≥ 2	13	-0.198** (0.068)	-0.251*** (0.047)
Mean offers	event year ≥ 0	23	-0.367 (0.839)	-0.116 (0.751)
	event year ≥ 2	12	-0.658 (0.504)	+0.119 (0.335)

ATT is the mean of the per-NSN τ_i , with HC1 (heteroskedasticity-robust) standard errors. *Main matching* picks the three nearest donors by Mahalanobis distance from the full eligible pool. *Same-FSG sensitivity* adds an exact-match constraint requiring each donor to share the treated NSN's two-digit Federal Supply Group (the first two digits of the four-digit FSC code).

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A23. DLA-Only: Matched-Controls Fit Diagnostics (Covariate Balance)

Match-quality statistic	Main analysis	DLA-Only
Mean SMD , before matching	1.386	1.369
Mean SMD , after matching	0.207	0.248
Max SMD , after matching	0.950	1.012
Share balanced ($ SMD < 0.1$), after matching	32%	32%

Standardized mean differences between each treated NSN and its matched control pool, averaged over treated NSNs and the four matching covariates (the propensity distance is excluded). Lower is better balance; $|SMD| < 0.1$ is the conventional threshold. Matching is re-run on DLA-only data here.

Table A24. DLA-Only: Synthetic Controls – Event-Time Path by Year (vs Main Analysis)

Outcome	Event year	DLA-Only N	DLA-Only mean	SE	Main N	Main mean
Domestic sourcing share	+0	21	-0.441***	0.064	23	-0.461***
	+1	18	-0.391***	0.070	20	-0.367***
	+2	13	-0.449***	0.101	15	-0.415***
	+3	3	-0.698**	0.104	3	-0.682**
Maximum logged unit price	+0	21	-0.092	0.069	22	-0.025
	+1	19	-0.072	0.044	20	-0.068
	+2	13	-0.289***	0.053	14	-0.291***
	+3	3	-0.158	0.179	3	-0.162
Mean offers	+0	21	+0.174	0.568	22	+0.153
	+1	18	-0.374	0.349	20	-0.311
	+2	13	-0.025	0.265	14	-0.051
	+3	3	+0.800**	0.174	3	+0.863

Aggregate synthetic-control effect (mean of per-NSN demeaned effects) at each post-waiver event year, for the DLA-Only sample beside the main-analysis aggregate, each with its own fit count N. SE and significance stars from a one-sample t-test on the per-NSN effects ($df = N - 1$); both mean columns carry stars.

Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table A25. DLA-Only: Synthetic Controls – Per-NSN ATT Distribution

Outcome	N fits	Median ATT	Mean ATT	SD across NSNs	Significant ($p < 0.05$)	Median pre-fit RMSPE	Median active donors / pool size
Domestic sourcing share	21	-0.382	-0.452	0.287	19 / 21	0.0033	61 / 13,905
Maximum logged unit price	21	-0.112	-0.078	0.277	1 / 21	0.0004	4,519 / 15,548
Mean offers	21	-0.015	+0.320	2.254	2 / 21	0.0071	1,927 / 2,979

ATTs are estimated per treated NSN using `synthdid` (Arkhangelsky et al. 2021). Significance is assessed against the random-subsample placebo distribution (200 placebos per fit). Pre-fit RMSPE is the root mean squared prediction error in the pre-waiver period – smaller means the synthetic tracks its treated NSN more closely before the waiver. The “active donors / pool size” column reports the median count of non-treated NSNs that receive non-trivial weight ($w > 1 / (2 \cdot \text{pool size})$) against the total eligible pool – lower ratios indicate that the synthetic relies on a small subset of the pool.